Doc-to-LoRA: Learning to Instantly Internalize Contexts

Rujikorn Charakorn [1]   **Edoardo Cetin** [1]   **Shinnosuke Uesaka** [* 2]   **Robert T. Lange** [1]

Abstract

Long input sequences are central to in-context learning, document understanding, and multi-step reasoning of Large Language Models (LLMs). However, the quadratic attention cost of Transformers makes inference memory-intensive and slow. While context distillation (`CD`) can transfer information into model parameters, per-prompt distillation is impractical due to training costs and latency. To address these limitations, we propose *Doc-to-LoRA* (`D2L`), a lightweight hypernetwork that meta-learns to perform approximate `CD` within a single forward pass. Given an unseen prompt, `D2L` generates a LoRA adapter for a target LLM, enabling subsequent queries to be answered without re-consuming the original context, reducing latency and KV-cache memory consumption during inference of the target LLM. On a long-context needle-in-a-haystack task, `D2L` successfully learns to map contexts into adapters that store the needle information, achieving near-perfect zero-shot accuracy at sequence lengths exceeding the target LLM's native context window by more than $4\times$. On real-world QA datasets with limited compute, `D2L` outperforms standard `CD` while significantly reducing peak memory consumption and update latency. We envision that `D2L` can facilitate rapid adaptation of LLMs, opening up the possibility of frequent knowledge updates and personalized chat behavior. Checkout our project at github.com/SakanaAI/doc-to-lora.

1. Introduction

LLMs are commonly adapted to documents, tasks, and user preferences by placing relevant information in the context window, also known as prompting or in-context learning (ICL, Brown et al., 2020). While effective and convenient, ICL is transient and memory-intensive; long prompts
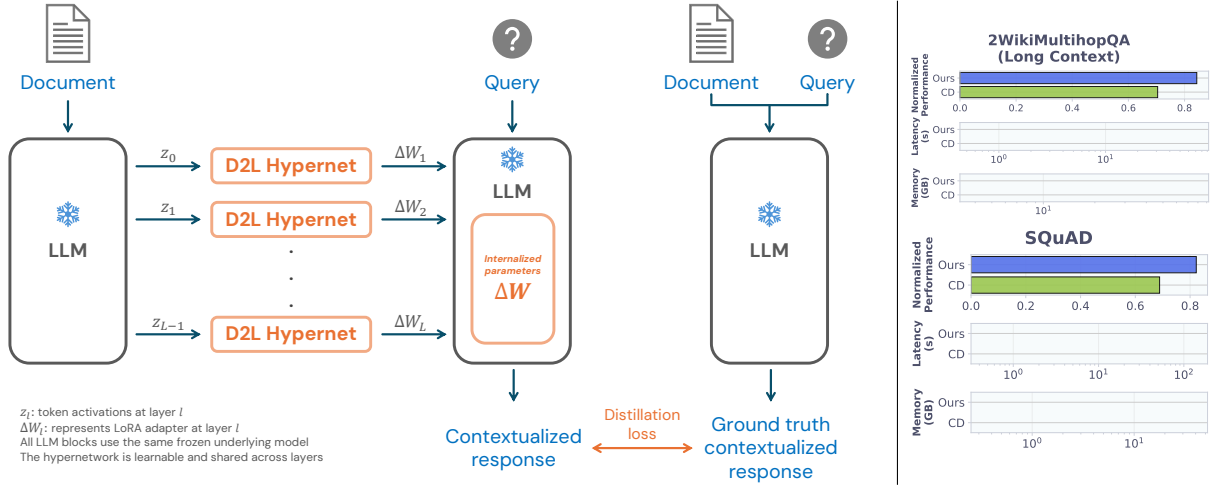
increase latency and memory via quadratic attention and KV-cache growth. Moreover, generation quality typically degrades under longer context lengths (Liu et al., 2024; Li et al., 2025; Hong et al., 2025). A standard remedy is to compress information from context into model parameters via supervised finetuning (SFT, Zhang et al., 2023; Pareja et al., 2025). However, SFT requires collecting a task-specific dataset that represents the desired behavior and risks overfitting when data is scarce. Furthermore, repeated expensive SFT processes are required if the information, e.g., user preference, changes over time. These constraints limit the possibility of scalable model customization by practitioners.

Context distillation (`CD`) offers a compelling alternative (Askell et al., 2021; Choi et al., 2023; Padmanabhan et al., 2023; Snell et al., 2023; Bhargava et al., 2024; Caccia et al., 2025; Eyuboglu et al., 2025; Kujanpää et al., 2025; Qi et al., 2025; Shin et al., 2025). At its core, `CD` trains an LLM—without access to relevant information—to imitate its own outputs when the information is provided in context. Once the information is *internalized*[1], subsequent inference calls are significantly faster as the model does not need the information to be in its context window. Still, its training is slow and memory-intensive, making it impractical for interactive or on-device applications. Furthermore, when the information sources are constantly changing, repeated slow `CD` processes must take place.

In this work, we aim to combine the convenience of ICL with the effective internalization of `CD`. We propose to meta-learn the `CD` process into a *hypernetwork* (Ha et al., 2016). In other words, we meta-train the hypernetwork to represent a mapping from a given context to its corresponding weight updates produced by `CD`. Conceptually, after reading a context once, a lightweight hypernetwork generates a context-specific LoRA adapter (Hu et al., 2022). With the adapter, the target LLM can respond to subsequent queries *without* the context, reducing latency and KV-cache size. Once trained, the hypernetwork can be reused for any new context. A single, inexpensive forward pass executes the learned `CD` process, enabling low-latency internalization. We call this method *Doc-to-LoRA* (`D2L`).

---
*S. Uesaka contributed while he was an intern at Sakana AI.
[1]Sakana AI, Tokyo, Japan [2]Minerva University, California, USA. Correspondence to: Rujikorn Charakorn <rujikorn@sakana.ai>.

---

[1]We define internalization concretely in Section 2. Informally, internalization is a process that transforms information into a model's parameters such that the model can later access it.

Figure 1. An overview of `D2L` training (left) and downstream performance (right). `D2L` learns to efficiently internalize information, outperforming traditional CD while significantly reducing latency and memory consumption across question-answering benchmarks under limited query budgets.

To enable `D2L` to internalize contexts longer than those seen during training, we base our hypernetwork on the Perceiver architecture (Jaegle et al., 2021), allowing it to map variable-length inputs to a fixed-shape LoRA adapter. Furthermore, `D2L` utilizes a chunking mechanism that allows it to produce higher-rank LoRA adapters for longer contexts without changing the output shape of the hypernetwork. This mechanism is crucial for processing documents longer than the context window of the target LLM. As a result, on a synthetic Needle-in-a-Haystack (NIAH) task, `D2L` achieves near-perfect zero-shot accuracy on contexts beyond 32K tokens, despite being trained only on sequences of up to 256 tokens, indicating that `D2L` is capable of generalizing beyond the lengths of its training data and that the chunking mechanism is effective at processing long sequences. Empirically, on real-world QA tasks, `D2L` learns an effective mapping between contexts and their corresponding LoRA adapters, outperforming `CD` under low compute budgets while significantly reducing internalization latency and compute cost. On long-document QA tasks, `D2L` generalizes in a zero-shot manner to documents exceeding the training length, without being explicitly trained to internalize such long inputs. A conceptual overview of `D2L` is shown in Figure 1. We summarize our contributions as follows:

- We introduce a meta-training objective that distills the `CD` process into a hypernetwork. This approach amortizes the overhead associated with traditional `CD` by training the hypernetwork to emulate the entire `CD` process in a single forward pass (Section 3.1).

- We introduce a well-designed architecture that makes the hypernetwork robust to varying input lengths and allows it to output higher-rank LoRAs by chunking long inputs (Section 3.2). Consequently, `D2L` can internalize needle information at near-perfect zero-shot accuracy up to $4\times$ the maximum length of the base

LLM on an NIAH task (Section 4).

- We empirically validate `D2L` and show that, under limited model update budgets, it outperforms `CD` with improved internalization efficiency, significantly reducing update latency and memory usage. Furthermore, `D2L` generalizes zero-shot to long-document QA tasks with even greater computational efficiency compared to the gains observed on shorter contexts (Section 5.1). Remarkably, `D2L` zero-shot transfers visual information effectively from a visual-language model (VLM) to a text-based LLM, allowing the target model to classify images purely through internalized information (Section 5.2). Finally, analyses of extreme generalization, performance gains, ablations, and failure modes of `D2L` are included in Section 6 and App. C.

2. Preliminaries

Context Distillation (`CD`) is a self-distillation method that internalizes behaviors induced by an in-context prompt, c, into a model's parameters. Unlike traditional knowledge distillation (Buciluă et al., 2006; Hinton et al., 2015), `CD` uses the same LLM for both the "teacher" and the "student." The teacher has access to the prompt c, while the student does not. Concretely, given a context and query pair (c, x), we sample a response y from the teacher, parameterized by θ: $y \sim p_\theta(\cdot \mid x, c)$. Then, y is used as the target for training the student with only access to the query x. Mathematically, the `CD` training objective can be written as

$$\min_{\theta_c} \left[\mathrm{KL}\big(p_\theta(y \mid x, c) \,\|\, p_{\theta_c}(y \mid x)\big) \right], \qquad (1)$$

where $\mathrm{KL}(\cdot \| \cdot)$ denotes the Kullback–Leibler divergence between two probability distributions and the context-specific parameters θ_c are initialized from the original parameters θ. Given that modern LLMs are effective at following in-

structions provided in context, `CD` is a practical way to internalize new knowledge or behaviors without expensive data gathering steps needed for SFT.

Defining Internalization. Optimizing Equation (1) risks overfitting since it only considers a single (c, x, y) triplet. We refer to it as *query-dependent* distillation as the learning signal comes solely from a single query x. A more robust alternative is *query-independent* distillation, which utilizes multiple queries, $X = \{x_i\}_{i=1}^n$, and responses, $Y = \{y_i \sim p_\theta(\cdot|x_i, c)\}_{i=1}^n$. Collectively, X and Y make a small dataset $\mathcal{D}_c = \{(x_i, y_i)\}_{i=1}^n$ specifically for internalizing context c with various queries and self-generated responses. Typically, X is generated by prompting an LLM to generate relevant questions for the context c (Caccia et al., 2025; Kujanpää et al., 2025; Eyuboglu et al., 2025). Thus, the query-independent distillation objective becomes

$$\min_{\theta_c} \quad \mathbb{E}_{(x,y)\sim\mathcal{D}_c}\Big[\mathrm{KL}\big(p_\theta(y \mid x, c) \parallel p_{\theta_c}(y \mid x)\big)\Big] \quad (2)$$

We define the optimization in Equation (2) as the *internalization* process of a context c into model parameters θ. Successful internalization means the model can access information from c through the internalized parameters θ_c, behaving as if c were provided in context. Context distillation has strong implications for real-world applications as it enables the creation of many specialized models without requiring explicit data collection. For instance, persistent instructions like safety and alignment prompts are prime candidates for internalization. These prompts are often tailored for different use cases and must remain in the context window throughout deployment.

3. Meta-Learning Context Distillation

3.1. Learning to Internalize Context with Hypernetworks

In this work, we focus on meta-learning query-independent `CD`, which allows the internalized parameters θ_c to be used with unseen downstream queries. Thus, the meta-training phase trains a hypernetwork to map a context c to a set of adapter parameters that modify a frozen base model θ, yielding a context-internalized model $\theta_c = \theta + \Delta W_c$, where $\Delta W_c = H_\phi(c)$ and H_ϕ represents a mapping from a raw string to LoRA parameters, parameterized by ϕ (Section 3.2). The overall meta-training process is similar to the vanilla `CD` (Eq. 2) with one key distinction: Optimizing a single H_ϕ that generalizes across many contexts (tasks), rather than optimizing separate ΔW per context. The hypernetwork must be exposed to a vast number of contexts to robustly represent the `CD` process. Consequently, the meta-training dataset $\mathcal{D}$ must include diverse contexts, queries, and responses: $\mathcal{D} = \{c_i, \mathcal{D}_{c_i}\}_{i=1}^n$. Mathematically, we optimize H_ϕ to minimize the divergence between the context-conditioned teacher and the context-internalized student:

$$\min_\phi \quad \mathbb{E}_{(c,\mathcal{D}_c)\sim\mathcal{D}} \; \mathbb{E}_{(x,y)\sim\mathcal{D}_c} J(x, y, c), \quad (3)$$

$$J(x, y, c) = \mathrm{KL}\big(p_\theta(y \mid x, c) \parallel p_{\theta + H_\phi(c)}(y \mid x)\big) \quad (4)$$

After training on a large corpus $\mathcal{D}$, the trained hypernetwork should serve as a generic mapping between context information c and corresponding internalized parameters θ_c, generated by a single `D2L` hypernetwork forward pass. That is, a trained mapping H amortizes *both* the query generation process and the backpropagation needed by traditional `CD`. The data generation pipeline is described in App. B.1.

3.2. `D2L` Architecture

Next, we describe how `D2L` maps a context string into LoRA matrices (Fig. 1 shows the overview of this process). A context c is fed through the frozen target LLM to obtain per-layer token activations $Z \in \mathbb{R}^{L \times N \times D}$, where L is the number of Transformer layers (including the embedding layer), N is the number of context tokens, and D is the hidden size. We denote the activations at depth l by $Z_l \in \mathbb{R}^{N \times D}$, with Z_0 being the token embeddings. For each transformer layer $l \in \{1, \ldots, L\}$, a shared hypernetwork h_ϕ consumes Z_{l-1} and outputs low-rank LoRA parameters (Hu et al., 2022): $h_\phi(Z_{l-1}) = \Delta W_l = B_l A_l$. Specifically, each target weight $W_l \in \mathbb{R}^{d_l^{\mathrm{out}} \times d_l^{\mathrm{in}}}$ is adapted as

$$W_l' = W_l + \alpha_l B_l A_l; \; A_l \in \mathbb{R}^{r \times d_l^{\mathrm{in}}}, \; B_l \in \mathbb{R}^{d_l^{\mathrm{out}} \times r}, \quad (5)$$

where r is the rank and α_l is a learnable per-layer scalar.

Perceiver-Based Hypernetwork: We structure the hypernetwork h_ϕ as a Perceiver-style cross-attention module (Jaegle et al., 2021) that maps variable-length inputs (Z_{l-1}) to a fixed number of latent queries (the LoRA rank). We briefly describe the simplest configuration of the Perceiver with a single cross-attention block here.[2] Specifically, for each layer l, we use r learnable, input-independent latent queries $Q_m \in \mathbb{R}^{r \times d_q}$. Cross-attending Q_m to Z_{l-1} yields r latent vectors:

$$U_l = \mathrm{XAttn}\big(Q_m, K(Z_{l-1}), V(Z_{l-1})\big) \in \mathbb{R}^{r \times d_u}, \quad (6)$$

where $K(\cdot)$ and $V(\cdot)$ are linear projections. Two per-layer output heads then produce the LoRA matrices by mapping each latent to a row of A_l and a column of B_l. Overall, the mapping from a context string to LoRA matrices, H_ϕ, given context c and its token activations Z, can be written as

$$H_\phi(c) = \Delta W = \{\Delta W_l\}_{l \in \{1, \ldots, L\}}, \quad (7)$$

$$\Delta W_l = B_l A_l = h_\phi(Z_{l-1}) \quad (8)$$

---

[2]The cross-attention operator can be repeated many times and interleaved with latent self-attention blocks.

Long-Context Composition via Chunking: For long contexts, we partition c into K contiguous chunks $\{c^{(k)}\}_{k=1}^K$ and process each chunk independently through the hypernetwork, producing per-chunk adapter $(A_l^{(k)}, B_l^{(k)})$. We then combine chunks by concatenating along the rank dimension

$$A_l = \begin{bmatrix} A_l^{(1)} \\ \vdots \\ A_l^{(K)} \end{bmatrix}, \qquad B_l = \begin{bmatrix} B_l^{(1)} & \cdots & B_l^{(K)} \end{bmatrix}, \quad (9)$$

resulting in LoRA matrices with the total rank $r \cdot K$. This composition allows `D2L` to integrate information across many chunks without changing the output shape of the hypernetwork. Note that the described architecture is not limited to generating LoRA matrices. In App. D, we empirically show that it can effectively learn to output compressed KV cache (i.e., prefix-tuning, Li & Liang, 2021).

4. Implanting Synthetic Needle-in-a-Haystack Information

In this section, we aim to show that `D2L` (i) successfully induces knowledge internalization, enabling the base model to recall the implanted information *without* reading the raw context, (ii) effectively bypasses the inherent context-length limitations of the base language model, and (iii) reduces the computational requirements for inference, especially when the inputs are long. For illustration purposes, we evaluate `D2L` on a synthetic needle-in-a-haystack (NIAH) information retrieval task.

Our needle-in-a-haystack (NIAH) task is primarily based on the RULER benchmark (Hsieh et al., 2024). Briefly, the NIAH task requires the model to locate a specific piece of information (needle) within a long, distracting document (haystack). The needle is a sentence defining a special 4-digit number, e.g., "`The special magic number is 0042.`" This sentence is randomly inserted into a document filled with distractor text. The goal is to accurately retrieve the number when prompted. We use `gemma-2-2b-it` (Team et al., 2024) with 8K context length as the base LLM in all experiments.

During `D2L`'s meta-training, we use input contexts ranging from 32 to 256 tokens in length. The training inputs are randomly chunked from 1 to 8 chunks with a minimum chunk size of 25 tokens. We use a simplified training setup for this pedagogical experiment (see App. A for more details).

During evaluation, the baseline has direct access to both the haystack and the query. For `D2L`, the base LLM does not have direct access to any part of the original context but is simply given the query prompt: "`What is the special magic`
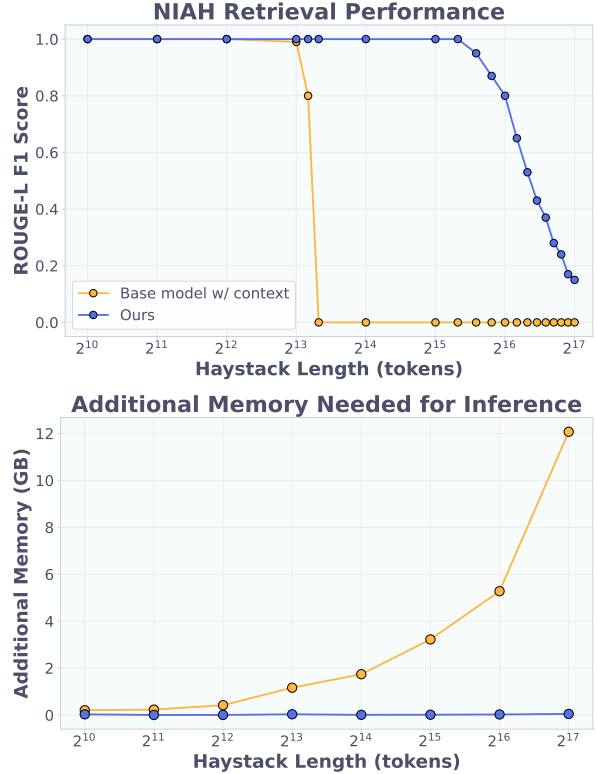

Figure 2. NIAH retrieval performance (top) and additional memory needed for inference (bottom).

`number? Reply with only the number.`" To perform well at this task, `D2L` must learn to map context information into a LoRA adapter that stores the value of the needle. With successful internalization, the adapted base model would be able to give a correct response based purely on the knowledge contained in the LoRA adapter. `D2L` processes the inputs by segmenting them into equal-sized chunks with 1024 tokens as the maximum chunk size. This chunk size is four times larger than the 256-token maximum sequence length seen during training.

`D2L` Learns Effective `CD` for NIAH Task and Generalizes Beyond Base Model's Context Length: The result in Figure 2 (left) confirms the effectiveness of our approach. `D2L` successfully internalizes the needle information and achieves perfect accuracy similar to the base model with in-context information up to 8K tokens. When the haystack is longer than 8K tokens, the base model's performance drops sharply due to its limited context length. In contrast, `D2L` maintains high retrieval accuracy across these longer sequences. Notably, performance remains close to perfect up to 40 chunks (40K tokens), which is *quintuple* the number of chunks the model has been exposed to during its training phase. Beyond this point, performance begins to degrade gracefully. These results demonstrate that `D2L` exhibits strong generalization capabilities for both chunk size and the total number of chunks.

D2L Reduces Inference Cost Under Long Inputs: The result in Figure 2 (right) shows that **D2L** not only achieves high accuracy but also demonstrates significant efficiency improvements, requiring less memory than the base model, particularly at extended context lengths. The base model uses more than 12 GB of additional memory to generate a response to a 128K-token haystack. In contrast, the model with internalized knowledge consistently uses significantly lower memory (< 50 MB) regardless of the length of the haystack. This result highlights a potential real-world application, where users first internalize long private documents, avoiding memory-intensive KV cache at inference.

5. Experiments

In this section, we move from the synthetic NIAH task to a more realistic setting, where **D2L** has to learn to approximate generic **CD** from a large corpus of English documents. We evaluate the ability of **D2L** to operate as a context internalizer for question-answering tasks on 6 real-world benchmarks, including short-passage and long-document QA tasks. A key advantage of **D2L** on these tasks is its ability to provide instant and inexpensive internalization, which allows the base LLM to answer subsequent queries without consuming the context window.

Experimental Setup: We use the data generated by the process described in App. B.1 to meta-train **D2L**. Our Perceiver-based **D2L** has 8 cross-attention blocks without self-attention layers. It splits the inputs into equal-sized chunks with 8K tokens and outputs a rank-8 LoRA adapter for each context chunk. Each generated adapter is applied to the "down projection" layer of each MLP block of the base model. In total, it has only 309M trainable parameters. During evaluation, **D2L** can operate in batched and iterative LoRA generation. In batched mode, it batches the token activations Z_l across all layers, producing LoRA adapters for all the layers within a single forward pass. Iterative mode, on the other hand, produces the adapter one layer at a time. The two modes allow us to prioritize either speed (batched) or lower memory consumption (iterative).[3]

We consider the following *in-parameter knowledge* baselines—where the model has to answer input queries based on its internal knowledge: (i) **CD (oracle)** serves as the empirical upper-bound of in-parameter knowledge methods as it optimizes the internalized weights directly on the target query (Eq. 1), (ii) **CD** trains the internalized weights using generated queries from gemma-3-12b-it (Eq. 2), (iii) **T2L** (Charakorn et al., 2025) is a hypernetwork-based method that directly maps the context to LoRA, similar to **D2L**. However, it has been trained on SFT datasets

with next-token prediction loss on ground-truth tokens, and (iv) **Base model (without context)** serves as the lower bound, as it does not have access to any context information. Both **CD** baselines internalize the context into a rank-8 LoRA adapter applied at the "down projection" layer. We use the publicly available checkpoint for T2L, which has been trained to output rank-8 adapters for the K and V projection layers. To better contextualize the reported performance, we also include *in-context knowledge* baselines—where the model can simply retrieve the answer directly from the provided context: (i) **Base model with context** serves as the upper-bound as it can directly access the answer in context, and (ii) **LLMLingua-2** (Pan et al., 2024), a prompt compression method. Results of other base models (Mistral-7B-Instruct-v0.2 and Qwen3-4B-Instruct-2507) are in App. E.

5.1. Question-Answering from Implanted Knowledge

The following experiments test **D2L**'s ability to accurately store and retrieve information effectively from generated LoRA adapters. We report the test performance on unseen instances and use word-level ROUGE-L F1 scores as the main task performance metric. The reported performance is relative to the base model with direct access to the contexts. We provide its absolute performance in Table 11.

5.1.1. EFFICIENT AND EFFECTIVE INTERNALIZATION ON READING COMPREHENSION TASKS

We assess performance on three standard reading comprehension benchmarks: **SQuAD** (span extraction, Rajpurkar et al., 2016), **DROP** (discrete reasoning over passages, Dua et al., 2019), and **ROPES** (reasoning with background knowledge, Lin et al., 2019). As shown in Figure 3 (left), **D2L** outperforms all the in-parameter baselines across all three benchmarks (see more results in App. C), achieving 82.5% relative performance compared to the ICL upper bound on the SQuAD benchmark. Compared to LLMLingua-2, **D2L** performs roughly similarly to compressing the context down to 40% of its original length, but it does so while removing the context entirely.

The main benefit of using **D2L** is its update efficiency. **D2L** not only outperforms the **CD** baseline given limited compute and time constraints, it significantly reduces update latency and memory (Figure 3, middle and right). Specifically, **D2L** instantly internalizes context information within less than a second using either batched or iterative mode. **CD** (oracle) uses around 40 seconds to internalize information due to repeated stochastic gradient descent updates. Vanilla **CD** requires even more time to internalize because of the additional query generation process, totaling more than 100 seconds. While T2L can instantly update the model, it cannot effectively internalize knowledge due to its highly

---

[3]Both modes are mathematically equivalent. Any task performance differences are caused by different low-level matmul kernels and the non-commutative nature of floating-point operations.
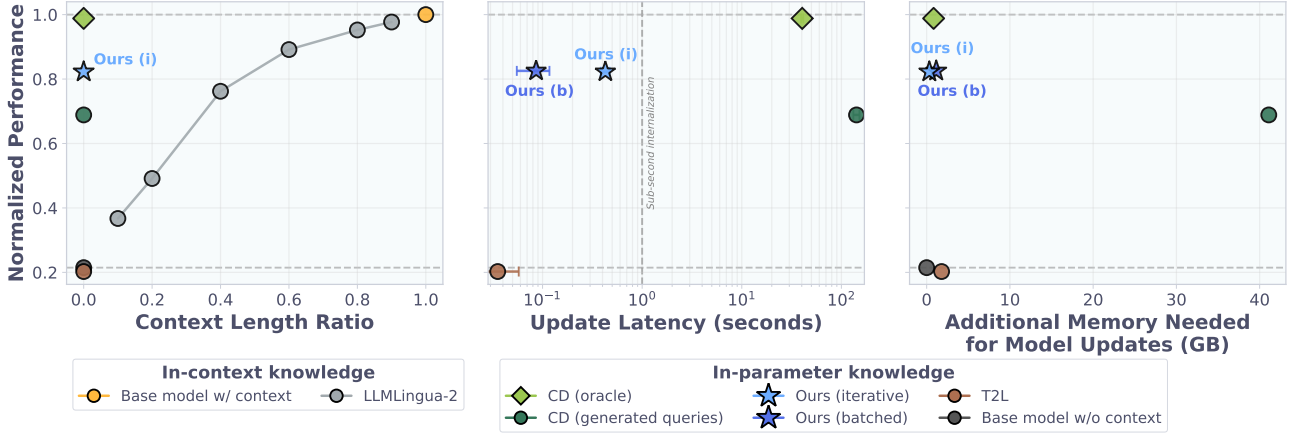
Figure 3. QA performance on SQuAD compared to the used context length ratio (left), update latency (middle), and additional memory needed for model updates (right). LLMLingua-2 compresses the input with [10%, 20%, 40%, 60%, 80%, 90%] compression rates from right to left (gray dots).

focused training data.

D2L requires low additional memory during the internalization process, similar to that of **CD** (oracle). Both use less than 2 GB of VRAM during the update process across the three benchmarks. In contrast, **CD** with generated queries uses more than 40 GB due to backpropagating on many queries at once. This result highlights the main practical benefit of **D2L**: providing effective instant internalization with low memory requirement. Due to the space limit, we provide results on DROP and ROPES benchmarks in App. C. The overall trends and observations are similar to the results on the SQuAD dataset shown in Figure 3. In the next experiment, the memory efficiency of **D2L** is much more impactful when considering longer context tasks.

5.1.2. INSTANT ZERO-SHOT INTERNALIZATION OF LONG-CONTEXT INFORMATION

We extend our evaluation to long-context scenarios, which pose a significant challenge for standard **CD** due to memory and computational constraints of distilling long context, using three document-grounded QA benchmarks from Long-Bench dataset (Bai et al., 2023), **2WikiMultihopQA** (Ho et al., 2020), **MultiFieldQA**, and **QASPER** (Dasigi et al., 2021). The length of test samples can go up to 32K tokens. We note that **D2L** has never seen such long sequences during training. Specifically, the longest training sample is 2,344 token long (see Figure 5). In this experiment, **CD** uses generated queries from truncated documents since the base model performs worse without truncation.

The results in Table 1 and Figure 4 show that **D2L** can effectively internalize long-context documents without being explicitly trained to do so. Like the previous experiment, **D2L** outperforms **CD** with generated queries and almost

reaches the upper-bound performance of the oracle **CD** on 2WikiMultihopQA. Even with 5 queries, **CD** uses up to 79 GB of VRAM to internalize the documents. Furthermore, the oracle **CD** requires more than 7 GB of VRAM during the update. In contrast, **D2L** with iterative LoRA generation uses 2x less memory compared to the oracle's update while maintaining sub-second internalization. Results on other benchmarks are included in App. C. Once again, this result underscores the practical benefits of **D2L** for real-time and on-device applications.

In this experiment, we also investigate the additional memory used *during response generation* of the base LLM since it requires a non-negligible amount of VRAM to read long-context information. Figure 4 shows that **D2L** answers more accurately than the **CD** baseline while significantly reduces memory used during response generation. Specifically, the ICL baseline requires around 1 GB of VRAM while all in-parameter knowledge methods require less than 100 MB.

Interestingly, on 2WikiMultihopQA and MultiFieldQA datasets, we observe that after internalizing a long context via **D2L**, the performance of the LLM slightly improves when reading the context directly again (indicated by `Ours + truncated context` in Figure 4). Also, when used with LLMLingua-2 with specific compression rates, the performance improves over feeding the base model with uncompressed truncated inputs. We hypothesize that this might be because of the *lost-in-the-middle* (Liu et al., 2024) or *attention noise* (Ye et al., 2025) phenomena. Intuitively, text-based compression reduces unnecessary words from the context, and, therefore, can reduce attention noise and improve performance. While **D2L** produces similar improvement, we believe that the improvement comes from different mechanisms. When the model faces strong attention noise or the answers are truncated, the LLM then

Table 1. Performance, update memory, and latency of in-parameter knowledge methods on 2WikiMultihopQA benchmark.

| Method	Normalized Performance ($\uparrow$)	Additional Update Memory (GB, $\downarrow$)	Mean Update Latency (s, $\downarrow$)
CD (oracle query)	0.901	7.820	40.171 ± 0.351
D2L (batched)	**0.857**	11.522	$\mathbf{0.209} \pm 0.123$
D2L (iterative)	0.844	**3.791**	0.551 ± 0.101
CD (25 generated queries)	0.745	59.925	465.454 ± 67.868
CD (5 generated queries)	0.704	79.371	72.537 ± 7.821

Figure 4. Long document QA performance. LLMLingua-2 compresses the input with $[20\%, 40\%, 60\%, 80\%, 90\%]$ compression rates from right to left (gray dots).

falls back on using its internal knowledge (Tao et al., 2024), which has been updated by D2L. We do not observe this effect in any of the benchmarks used in the previous experiment, presumably because attention noise does not occur under shorter context lengths.

5.2. Zero-Shot Internalization of Visual Information

Table 2. Using a VLM as the context encoder.

	SQuAD	DROP	ROPES	Imagenette
D2L (LLM $\Rightarrow$ LLM)	**0.814**	**0.655**	**0.906**	N/A
D2L (VLM $\Rightarrow$ LLM)	0.705	0.568	0.772	**75.03%**

Since D2L is based on the Perceiver architecture that can map arbitrary-length inputs to fixed-size outputs, we are not restricted to using the target LLM as the context encoder. In this experiment, we explore the feasibility of using the proposed architecture for bridging between a VLM (gemma-3-4b-it) and a pure-text model (gemma-2-2b-it). We want to see whether D2L, *without seeing any images during training*, can internalize visual information from the VLM zero-shot. Since the number of layers of the VLM and the LLM are not the same, D2L maps the activations from the first 26 layers of the VLM to the corresponding layers of the target LLM. We keep the rest of the training setup identical to the main experiment. Only the language modeling part of the VLM is used during training. To test whether the target model can classify images based solely on the internalized information, we use the Imagenette dataset (Howard, 2019), a 10-class subset of ImageNet (Russakovsky et al., 2015). The target

model is prompted with the following text: "`What is in this image? Choose exactly one of the following classes: tench, English springer, cassette player, chainsaw, church, French horn, garbage truck, gas pump, golf ball, parachute. Respond with only the correct class without any other text.`" Achieving better-than-random accuracy (10%) on this task means that D2L can zero-shot transfer visual information to a text-only LLM.

The results are shown in Table 2. D2L successfully learns to map the VLM's activations into LoRA matrices of the target LLM. Although using a VLM as the context encoder negatively impacts text-based QA performance, it enables us to directly communicate visual information extracted by the VLM into the parameters of the target text-only model. Specifically, the target model gets 75.03% accuracy purely through the internalized information. This is a striking result given that D2L and the target LLM have never seen information from any other modalities except text.

6. Analyses

Zero-Shot Query Internalization: To further test the robustness of D2L, we use the existing SQuAD dataset with a slight modification. Instead of internalizing the document, we use D2L to internalize the query. This means that, during evaluation, the target model sees the document directly but not the query. This experiment represents an extreme generalization test for D2L, as it is trained to internalize

Table 3. SQuAD query internalization.

Method	Recall	Precision
Base model w/ context	0.886	0.876
D2L	0.740	0.720
D2L (swapped)	0.587	0.044
Base model w/o context	0.185	0.205

Table 4. SQuAD (100 samples) performance with varying queries.

Method	Normalized Performance	Update Latency (s)
CD (oracle)	0.988	8.763 ± 0.313
D2L	**0.866**	**0.086** ± 0.061
CD (100 generated queries)	0.650	631.101 ± 12.879
CD (50 generated queries)	0.601	311.661 ± 7.211
CD (20 generated queries)	0.506	129.044 ± 5.813

knowledge, not queries. The results in Table 3 are encouraging and suggest that D2L can function properly under such an extreme generalization test. In this "swapped" configuration, while D2L performs worse, it still achieves moderate performance on the ROUGE-L recall metric (0.587) and outperforms the no-context baseline (0.185). We observe that the query-internalized model occasionally generates correct answers, but its outputs are typically verbose, leading to a sharp decrease in precision. This test suggests that D2L could be used to internalize other kinds of information beyond factual information from documents. Qualitative examples of this behavior are provided in Figure 8.

D2L Emulates CD over Many Generated Queries: In the main experiments (Section 5), D2L consistently outperforms vanilla CD under small query budgets. We hypothesize that D2L performs better because the hypernetwork learns to emulate the effect of distilling over a large number of queries per context. Although each training example provides only 10 generated queries, training across millions of context samples exposes the hypernetwork to a larger and more diverse query distribution. In other words, the training samples collectively regularize D2L to be robust to a wider variety of queries than presented in any single sample.

Empirically, increasing the number of generated queries for vanilla CD improves its performance, but it remains well below that of D2L (Table 4). On a 100-sample subset of SQuAD, its performance rises from 0.506 (20 queries) to 0.650 (100 queries). Crucially, CD with 100 queries takes more than 10 minutes to internalize *each sample*. In contrast, D2L achieves 0.866, substantially closer to the oracle CD upper bound, without incurring per-sample query generation or backpropagation. Although increasing the budget to hundreds or thousands of queries would likely boost CD performance beyond that of D2L, the update latency would also increase substantially. Such a large latency would render the user experience non-responsive, especially in fast-changing environments, e.g., active codebases or agentic tasks. This result highlights the effectiveness of D2L at instantly inter-

nalizing knowledge in the sub-second regime. This trend supports the interpretation that the hypernetwork has learned a mapping that approximates the CD process with many more queries per context. A complementary explanation is that D2L has learned a specialized but biased form of CD. Specifically, D2L might assume that subsequent queries will always be related to the internalized knowledge (Table 8). D2L's training curves and ablation studies on training data sources and LoRA rank can be found in App. C.

7. Related Work

Hypernetworks have been used to adapt LLMs on-the-fly for different use cases (Ivison & Peters, 2022; Ivison et al., 2023; Phang et al., 2023; Lv et al., 2024; Charakorn et al., 2025). Notably, MEND (Li et al., 2024) trains a hypernetwork via the CD objective to compress few-shot examples into prefix tokens (Li & Liang, 2021). Similarly, Gisting (Mu et al., 2024) uses the CD objective to train "gist" tokens that compress task instructions. Cartridges (Eyuboglu et al., 2025) aims to utilize abundant *sleep-time* compute budget for CD based on the prefix-tuning parameterization. In contrast to prior works, D2L aims to capture a generic CD process that can be applied to arbitrary information presented as context while being time- and memory-efficient.

A closely related method, Generative Adapter (GA, Chen et al., 2025), optimizes a hypernetwork using the next-token prediction loss on ground-truth tokens. GA first trains its hypernetwork on a pre-training corpus and then finetunes it on curated queries and responses from existing SFT datasets. In contrast, D2L meta-trains the hypernetwork with the CD objective—using primarily generated queries and self-responses—to output a context-specific LoRA adapter. Empirical results in App. C.2 also suggest that using the CD objective is crucial for robust generalization of the hypernetwork. We also note that GA achieves a significantly higher ROUGE-L F1 score than the base model (see Table 12) because of its training method that directly trains the model to output the ground truth tokens via the SFT loss. However, it achieves much lower ROUGE-L recall, indicating that the improved F1 scores comes solely from the fact that Generative Adapter outputs much shorter responses despite answering with less factual accuracy (represented by the ROUGE-L recall score). Finally, the use of generated queries and self-responses allows straightforward extensions of D2L to domains where finetuning datasets are not available. Due to the space constraint, extended Related Work and Conclusion sections are provided in App. F and G.

Acknowledgement

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

References

Askell, A., Bai, Y., Chen, A., Drain, D., Ganguli, D., Henighan, T., Jones, A., Joseph, N., Mann, B., DasSarma, N., et al. A general language assistant as a laboratory for alignment. *arXiv preprint arXiv:2112.00861*, 2021.

Bai, Y., Lv, X., Zhang, J., Lyu, H., Tang, J., Huang, Z., Du, Z., Liu, X., Zeng, A., Hou, L., Dong, Y., Tang, J., and Li, J. Longbench: A bilingual, multitask benchmark for long context understanding, 2023.

Bhargava, A., Witkowski, C., Detkov, A., and Thomson, M. Prompt baking. *arXiv preprint arXiv:2409.13697*, 2024.

Bourtoule, L., Chandrasekaran, V., Choquette-Choo, C. A., Jia, H., Travers, A., Zhang, B., Lie, D., and Papernot, N. Machine unlearning. In *2021 IEEE symposium on security and privacy (SP)*, pp. 141–159. IEEE, 2021.

Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D., Wu, J., Winter, C., Hesse, C., Chen, M., Sigler, E., Litwin, M., Gray, S., Chess, B., Clark, J., Berner, C., McCandlish, S., Radford, A., Sutskever, I., and Amodei, D. Language models are few-shot learners. In Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M., and Lin, H. (eds.), *Advances in Neural Information Processing Systems*, volume 33, pp. 1877–1901. Curran Associates, Inc., 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf.

Buciluǎ, C., Caruana, R., and Niculescu-Mizil, A. Model compression. In *Proceedings of the 12th ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 535–541, 2006.

Caccia, L., Ansell, A., Ponti, E., Vulić, I., and Sordoni, A. Training plug-and-play knowledge modules with deep context distillation. In *Second Conference on Language Modeling*, 2025. URL https://openreview.net/forum?id=ghyyHZYORi.

Cao, B., Cai, D., and Lam, W. Infiniteicl: Breaking the limit of context window size via long short-term memory transformation. *arXiv preprint arXiv:2504.01707*, 2025.

Charakorn, R., Cetin, E., Tang, Y., and Lange, R. T. Text-to-loRA: Instant transformer adaption. In *Forty-second International Conference on Machine Learning*, 2025. URL https://openreview.net/forum?id=zWskCdu3QA.

Chen, T., Fang, H., Xia, P., Liu, X., Durme, B. V., Zettlemoyer, L., Gao, J., and Cheng, H. Generative adapter: Contextualizing language models in parameters with a single forward pass. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=bc3sUsS6ck.

Chevalier, A., Wettig, A., Ajith, A., and Chen, D. Adapting language models to compress contexts. In Bouamor, H., Pino, J., and Bali, K. (eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 3829–3846, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.232. URL https://aclanthology.org/2023.emnlp-main.232.

Choi, E., Jo, Y., Jang, J., Jang, J., and Seo, M. Fixed input parameterization for efficient prompting. In *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 8428–8441, 2023.

Dasigi, P., Lo, K., Beltagy, I., Cohan, A., Smith, N. A., and Gardner, M. A dataset of information-seeking questions and answers anchored in research papers. 2021.

Dua, D., Wang, Y., Dasigi, P., Stanovsky, G., Singh, S., and Gardner, M. DROP: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proc. of NAACL*, 2019.

Eyuboglu, S., Ehrlich, R., Arora, S., Guha, N., Zinsley, D., Liu, E., Tennien, W., Rudra, A., Zou, J., Mirhoseini, A., et al. Cartridges: Lightweight and general-purpose long context representations via self-study. *arXiv preprint arXiv:2506.06266*, 2025.

Fang, J., Jiang, H., Wang, K., Ma, Y., Shi, J., Wang, X., He, X., and Chua, T.-S. Alphaedit: Null-space constrained model editing for language models. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=HvSytvg3Jh.

Ha, D., Dai, A., and Le, Q. V. Hypernetworks. *arXiv preprint arXiv:1609.09106*, 2016.

Hinton, G., Vinyals, O., and Dean, J. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

Ho, X., Duong Nguyen, A.-K., Sugawara, S., and Aizawa, A. Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps. In *Proceedings of the 28th International Conference on Computational Linguistics*, pp. 6609–6625, Barcelona, Spain (Online), December 2020. International Committee on Computational Linguistics. URL https://www.aclweb.org/anthology/2020.coling-main.580.

Hong, K., Troynikov, A., and Huber, J. Context rot: How increasing input tokens impacts llm performance. Technical report, Chroma, July 2025. URL https://research.trychroma.com/context-rot.

Howard, J. Imagenette: A smaller subset of 10 easily classified classes from imagenet, March 2019. URL https://github.com/fastai/imagenette.

Hsieh, C.-P., Sun, S., Kriman, S., Acharya, S., Rekesh, D., Jia, F., and Ginsburg, B. RULER: What's the real context size of your long-context language models? In *First Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=kIoBbc76Sy.

Hu, E. J., yelong shen, Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=nZeVKeeFYf9.

Ivison, H. and Peters, M. E. Hyperdecoders: Instance-specific decoders for multi-task nlp. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pp. 1715–1730, 2022.

Ivison, H., Bhagia, A., Wang, Y., Hajishirzi, H., and Peters, M. E. Hint: Hypernetwork instruction tuning for efficient zero-and few-shot generalisation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11272–11288, 2023.

Jaegle, A., Gimeno, F., Brock, A., Vinyals, O., Zisserman, A., and Carreira, J. Perceiver: General perception with iterative attention. In *International conference on machine learning*, pp. 4651–4664. PMLR, 2021.

Kujanpää, K., Marttinen, P., Valpola, H., and Ilin, A. Efficient knowledge injection in LLMs via self-distillation. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856. URL https://openreview.net/forum?id=drYpdSnRJk.

Li, T., Zhang, G., Do, Q. D., Yue, X., and Chen, W. Long-context LLMs struggle with long in-context learning. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856. URL https://openreview.net/forum?id=Cw2xlg0e46.

Li, X. L. and Liang, P. Prefix-tuning: Optimizing continuous prompts for generation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 4582–4597, 2021.

Li, Y., Ma, X., Lu, S., Lee, K., Liu, X., and Guo, C. MEND: Meta demonstration distillation for efficient and effective in-context learning. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=2Y5kBPtU0o.

Lin, J., Zettlemoyer, L., Ghosh, G., Yih, W.-T., Markosyan, A., Berges, V.-P., and Oğuz, B. Continual learning via sparse memory finetuning. *arXiv preprint arXiv:2510.15103*, 2025.

Lin, K., Tafjord, O., Clark, P., and Gardner, M. Reasoning over paragraph effects in situations. In *MRQA@EMNLP*, 2019.

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. doi: 10.1162/tacl_a_00638. URL https://aclanthology.org/2024.tacl-1.9/.

Lozhkov, A., Ben Allal, L., von Werra, L., and Wolf, T. Fineweb-edu: the finest collection of educational content, 2024. URL https://huggingface.co/datasets/HuggingFaceFW/fineweb-edu.

Lv, C., Li, L., Zhang, S., Chen, G., Qi, F., Zhang, N., and Zheng, H.-T. HyperloRA: Efficient cross-task generalization via constrained low-rank adapters generation. In *ACL Findings 2024*, 2024. URL https://openreview.net/forum?id=xa4GYUSvhW.

Mu, J., Li, X., and Goodman, N. Learning to compress prompts with gist tokens. *Advances in Neural Information Processing Systems*, 36, 2024.

Olah, C., Turner, N. L., and Conerly, T. A toy model of interference weights, 2025. URL https://transformer-circuits.pub/2025/interference-weights/index.html.

Padmanabhan, S., Onoe, Y., Zhang, M., Durrett, G., and Choi, E. Propagating knowledge updates to lms through distillation. *Advances in Neural Information Processing Systems*, 36:47124–47142, 2023.

Pan, Z., Wu, Q., Jiang, H., Xia, M., Luo, X., Zhang, J., Lin, Q., Ruhle, V., Yang, Y., Lin, C.-Y., Zhao, H. V., Qiu, L., and Zhang, D. LLMLingua-2: Data distillation for efficient and faithful task-agnostic prompt compression.

In Ku, L.-W., Martins, A., and Srikumar, V. (eds.), *Findings of the Association for Computational Linguistics ACL 2024*, pp. 963–981, Bangkok, Thailand and virtual meeting, August 2024. Association for Computational Linguistics. URL https://aclanthology.org/2024.findings-acl.57.

Pareja, A., Nayak, N. S., Wang, H., Killamsetty, K., Sudalairaj, S., Zhao, W., Han, S., Bhandwaldar, A., Xu, G., Xu, K., Han, L., Inglis, L., and Srivastava, A. Unveiling the secret recipe: A guide for supervised fine-tuning small LLMs. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=eENHKMTOfW.

Phang, J., Mao, Y., He, P., and Chen, W. Hypertuning: Toward adapting large language models without backpropagation. In *International Conference on Machine Learning*, pp. 27854–27875. PMLR, 2023.

Qi, S., Yang, B., Jiang, K., Wang, X., Li, J., Zhong, Y., Yang, Y., and Zheng, Z. In-context editing: Learning knowledge from self-induced distributions. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=w6rHCuN3YG.

Rajpurkar, P., Zhang, J., Lopyrev, K., and Liang, P. SQuAD: 100,000+ questions for machine comprehension of text. In Su, J., Duh, K., and Carreras, X. (eds.), *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 2383–2392, Austin, Texas, November 2016. Association for Computational Linguistics. doi: 10.18653/v1/D16-1264. URL https://aclanthology.org/D16-1264.

Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M., Berg, A. C., and Fei-Fei, L. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015. doi: 10.1007/s11263-015-0816-y.

Shi, H., Xu, Z., Wang, H., Qin, W., Wang, W., Wang, Y., Wang, Z., Ebrahimi, S., and Wang, H. Continual learning of large language models: A comprehensive survey. *ACM Comput. Surv.*, May 2025. ISSN 0360-0300. doi: 10.1145/3735633. URL https://doi.org/10.1145/3735633. Just Accepted.

Shin, H., Ji, L., Gong, Y., Kim, S., Choi, E., and Seo, M. Generative prompt internalization. In Chiruzzo, L., Ritter, A., and Wang, L. (eds.), *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 7338–7363, Albuquerque, New Mexico, April 2025. Associa-

tion for Computational Linguistics. ISBN 979-8-89176-189-6. URL https://aclanthology.org/2025.naacl-long.376/.

Snell, C. V., Klein, D., and Zhong, R. Learning by distilling context, 2023. URL https://openreview.net/forum?id=am22IukDiKf.

Su, J., Ahmed, M., Lu, Y., Pan, S., Bo, W., and Liu, Y. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.

Tao, Y., Hiatt, A., Haake, E., Jetter, A., and Agrawal, A. When context leads but parametric memory follows in large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 4034–4058, 2024.

Team, G., Riviere, M., Pathak, S., Sessa, P. G., Hardin, C., Bhupatiraju, S., Hussenot, L., Mesnard, T., Shahriari, B., Ramé, A., et al. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.

Team, G., Kamath, A., Ferret, J., Pathak, S., Vieillard, N., Merhej, R., Perrin, S., Matejovicova, T., Ramé, A., Rivière, M., et al. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025.

von Oswald, J., Henning, C., Grewe, B. F., and Sacramento, J. Continual learning with hypernetworks. In *International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=SJgwNerKvB.

Wang, Y., Ma, D., and Cai, D. With greater text comes greater necessity: Inference-time training helps long text generation. In *First Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=dj9x6JuiD5.

Ye, T., Dong, L., Xia, Y., Sun, Y., Zhu, Y., Huang, G., and Wei, F. Differential transformer. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=OvoCm1gGhN.

Zhang, P., Liu, Z., Xiao, S., Shao, N., Ye, Q., and Dou, Z. Long context compression with activation beacon. In *The Thirteenth International Conference on Learning Representations*, 2025a. URL https://openreview.net/forum?id=1eQT9OzfNQ.

Zhang, S., Dong, L., Li, X., Zhang, S., Sun, X., Wang, S., Li, J., Hu, R., Zhang, T., Wu, F., et al. Instruction tuning for large language models: A survey. *arXiv preprint arXiv:2308.10792*, 2023.

Zhang, Z., Rossi, R. A., Kveton, B., Shao, Y., Yang, D., Zamani, H., Dernoncourt, F., Barrow, J., Yu, T., Kim, S., Zhang, R., Gu, J., Derr, T., Chen, H., Wu, J., Chen, X., Wang, Z., Mitra, S., Lipka, N., Ahmed, N. K., and Wang, Y. Personalization of large language models: A survey. *Transactions on Machine Learning Research*, 2025b. ISSN 2835-8856. URL https://openreview.net/forum?id=tf6A9EYMo6. Survey Certification.

Zhao, D., Kobayashi, S., Sacramento, J., and von Oswald, J. Meta-learning via hypernetworks. In *4th Workshop on Meta-Learning at NeurIPS 2020 (MetaLearn 2020)*. NeurIPS, 2020.

A. NIAH experiment details

Here, we provide more details of the NIAH experiment. We use the cross-entropy loss on the ground-truth tokens instead of the self-distillation objective for ease of experimentation. Additionally, we use a simplified architecture where `D2L` maps the token activations from the 6th layer of `gemma-2-2b-it` to LoRA of all layers. The per-layer output heads in this architecture share the same inputs. The query used for training is the same as the evaluation query. Each training input is randomly chunked with the following probabilities: 50% for 1 chunk, 12% for 2 chunks, 37.5% for 3-8 chunks with equal chances. The training samples are 32- to 256-token long. There is a total of 640K training samples. `D2L` is trained for 1 epoch with learning rate 4×10^{-5}. The meta-training takes around 3 hours on a single H200 GPU. All evaluations and measurements are done with a single H200 GPU.
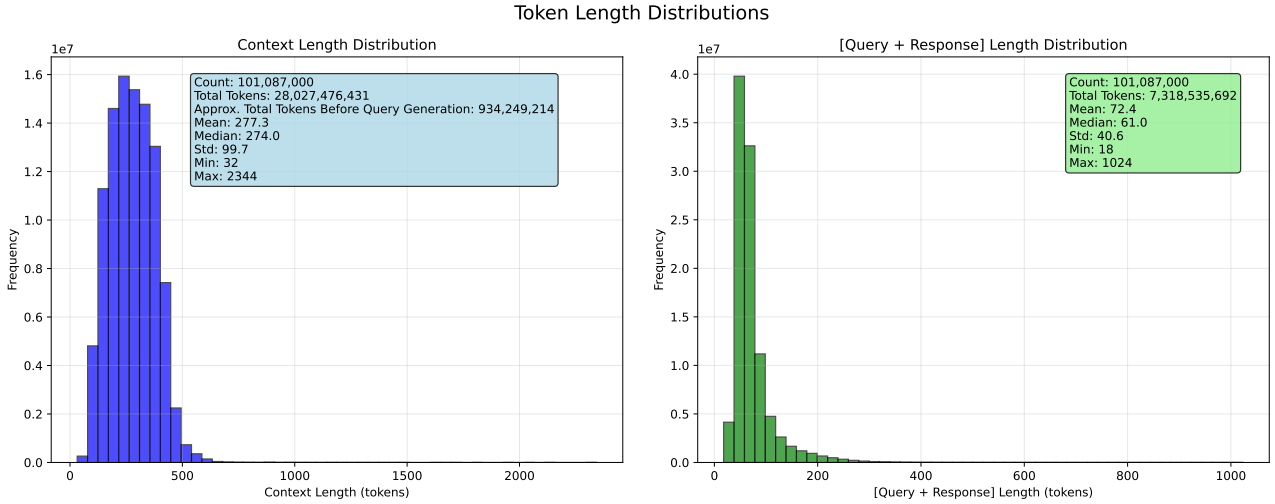
B. Main Experiments Details

Figure 5. Training data length distribution. `D2L` takes contexts while the base model takes both queries and responses to compute the loss during meta-training. The total count represents the total number of unique context-query-response triplets. The number of tokens from the original contexts before query generation is roughly around 900M tokens.

B.1. Meta-Training Data Generation Pipeline

We construct the meta-training dataset from a subset of FineWeb-Edu (Lozhkov et al., 2024), treating each sample as a context c. The subset contains approximately 900 million tokens. We further include passage-grounded QA datasets—PwC (Chevalier et al., 2023), SQuAD (Rajpurkar et al., 2016), ROPES (Lin et al., 2019), and DROP (Dua et al., 2019). After filtering out passages longer than $10,000$ characters, the combined corpus comprises approximately 3.2 million unique contexts. For FineWeb-Edu contexts, we generate 10 context-grounded queries per sample using `gemma-3-12b-it` (Team et al., 2025). We use a larger model for generating queries, ensuring better quality control and that the generated queries are mostly grounded in context. In practice, one could simply use the base model for generating queries. We prompt the model in two iterations, producing five queries per iteration. The first iteration uses Listing 4. The second uses Listing 5, which includes all previously generated query–answer pairs as in-context examples to encourage non-overlapping and increasingly challenging queries. The generated answers are discarded and not used for training. We augment the generated queries by putting each sample into a template instruction randomly chosen from Listing 6. We do the same for other datasets except PwC. For each unique context–query pair, we sample a single response from `gemma-2-2b-it` using Listing 7 and record the top-16 token logit values for every generated token. The logit values are the training target used in Equation (4). The overall training data length distribution is shown in Figure 5.

B.2. Training and Experimental Details

All evaluations and measurements are done with a single H200 GPU. We set the maximum number of test samples in each dataset to be 500 due to the high overhead latency of `CD`. We use the same query template for all the benchmarks except for

QASPER (see Listing 8).

The hypernetwork consists of two modules: a Perceiver-style cross-attention encoder that consumes per-layer token activations (Listing 1) and output heads that map the latent queries to LoRA matrices Listing 2). A pseudocode for the forward pass of the internalization process is shown in Listing 3.

We observe that the training can be unstable if `D2L` is trained to output multiple LoRAs from the beginning. Thus, we train `D2L` in a two-stage learning setup. First, `D2L` is trained to always output only one chunk for each input context for 80K gradient steps. This stage trains `D2L` to purely internalize information without emphasis on the compositionality of the generated LoRAs. Then, similar to the NIAH experiment, we randomly chunk each input with the following chunking probabilities: 50% for 1 chunk, 12% for 2 chunks, 37.5% for 3-8 chunks with equal chances. `D2L` is trained for 20K steps under the chunking setup. This stage regularizes `D2L` to output composable LoRAs. Both stages use the same data distribution shown in Figure 5. During training, each batch packs the context inputs into a 4K-token sequence and uses gradient accumulation, totaling more than 200K context tokens across 8 GPUs. We observe that this is crucial for good performance. Otherwise, the training converges too early with smaller batch sizes.
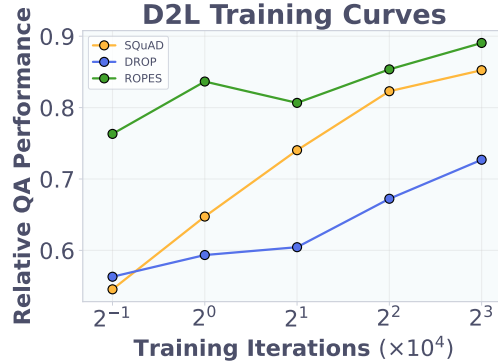
C. Additional Results (Gemma-2-2B)

Figure 6. Relative QA performance from various checkpoints throughout the training process.

C.1. Data Ablation

In previous experiments, training samples from QA tasks (SQuAD, DROP, ROPES) are included during training to ensure that `D2L` learns to respond correctly to the question formats presented in those datasets. Here, we remove samples from the QA tasks and exclusively use the training data derived solely from the pre-training FineWeb-Edu corpus.

Table 5. Data ablation

	SQuAD	DROP	ROPES
`D2L`	0.814	**0.655**	0.906
`D2L` w/o QA data	**0.838**	0.574	**0.923**
`CD` (20 queries)	0.689	0.504	0.776

Table 5 shows that the overall performance of `D2L` without QA data is comparable to the default training setup. Notably, it slightly outperforms the default `D2L` on SQuAD and ROPES, while underperforming significantly on DROP due to differences in training data distributions. Regardless of the QA data, D2L consistently outperforms low-budget CD (20 queries) across datasets, suggesting that D2L is largely robust and not highly sensitive to the training data format.

C.2. Ablating Training Objective

This experiment compares an alternative training loss used to train the hypernetwork. By default, `D2L` minimizes the context-conditioned KL objective in Equation (4) using self-generated responses from the base LLM as distillation targets (self-responses). We compare this default approach against an alternative that replaces KL with the next-token prediction (NTP) loss. Due to the high cost of training the

Table 6. Training loss and data ablation.

Method	SQuAD (normalized F1)	SQuAD (swapped, recall)
`D2L` (50%, KL)	**0.819**	**0.385**
`D2L` (50%, NTP)	0.763	0.235

hypernetwork, we use an early checkpoint (50%) for each method in this experiment. As shown in Table 6, the KL variant outperforms the NTP variant on SQuAD, achieving a normalized F1 score of 0.819 compared to 0.763. The gap is even wider when considering the "swapped" configuration—an extreme generalization case—where the KL approach achieves

a 0.385 recall score while the NTP variant only reaches 0.235. A plausible explanation is that KL distillation transfers richer information by matching the full teacher distribution $p_\theta(\cdot|x, c)$, preserving the uncertainty and alternative modes of the base LLM. Therefore, KL could propagate knowledge more effectively than NTP. This finding aligns with prior work (Padmanabhan et al., 2023; Eyuboglu et al., 2025; Caccia et al., 2025).

C.3. Increasing LoRA Rank

Effectively utilizing more training compute is important for improving the performance of modern machine learning models. In this experiment, we train D2L with more compute by increasing the rank of the generated LoRA to 16 (from 8) to probe its scalability. Table 7 shows that D2L with a higher rank (16) benefits from increased training compute. We observe that rank-16 LoRA outperforms rank-8 LoRA

Table 7. LoRA rank ablation

	SQuAD	DROP	ROPES
`D2L` (rank-8)	0.814	0.655	**0.906**
`D2L` (rank-16)	**0.896**	**0.711**	0.895

by significant margins on two benchmarks (SQuAD and DROP). Both perform similarly on the ROPES benchmark. These results suggest that increasing the capacity of LoRA can enhance performance and demonstrate that D2L can benefit from higher-capacity parameterizations. We opt to use rank-8 LoRA for the main experiments for ease of running experiments (less VRAM and faster training).

C.4. Knowledge Interference

We investigate further whether `D2L` can cause *knowledge interference*, where the internalized knowledge overrides existing internal knowledge in the base LLM. `D2L` is not trained to retain existing knowledge of the target model. Thus, the goal of this experiment is to observe the behavior of `D2L` without a dedicated knowledge preservation mechanism and set a baseline for how our simple training pipeline induces knowledge interference. To study this phenomenon, we replace the context of each sample in the SQuAD dataset with either an *assistant* prompt or a *distracting* prompt. The assistant prompt is simply the

Table 8. SQuAD with replaced contexts.

Method	SQuAD (assistant)	SQuAD (distracting)
Base model w/ replaced context	0.201	0.175
`D2L`	0.096	0.126
`CD` (10 generated queries)	0.211	0.203

text ``You are a useful assistant.`` while the distracting prompt is a book chapter from Project Gutenberg ($\sim$ 4K tokens). This dataset allows us to test `D2L` when the internalized knowledge and the queries are unrelated. Table 8 shows that `D2L` significantly reduces the performance compared to the base model and `CD`. We hypothesize that `D2L` might acquire a strong prior, assuming that the subsequent queries will always be related to the internalized knowledge. This might be caused by the bias in the data generation pipeline that only considers queries related to the input context. We think that by adding irrelevant queries into the training data, one could regularize the hypernetwork and help mitigate this bias. Other dedicated continual learning mechanisms, such as constraint edits (Fang et al., 2025) and sparse memory finetuning (Lin et al., 2025), could be incorporated to directly tackle the knowledge interference problem in future work.

Table 9. Performance, update memory, and latency of in-parameter knowledge methods on MultiFieldQA benchmark.

Method	Rel. Perf vs Truncated ICL ($\uparrow$)	Peak Update Memory (GB, $\downarrow$)	Mean Update Latency (s, $\downarrow$)
`CD` (oracle query)	1.041	7.820	41.912 ± 2.257
`D2L` (batched)	0.481	19.421	0.198 ± 0.094
`D2L` (iterative)	0.485	3.675	0.522 ± 0.098
`CD` (25 generated queries, mini-batch SGD)	0.528	53.232	431.479 ± 70.916
`CD` (5 generated queries, full-batch SGD)	0.419	40.231	81.406 ± 7.967
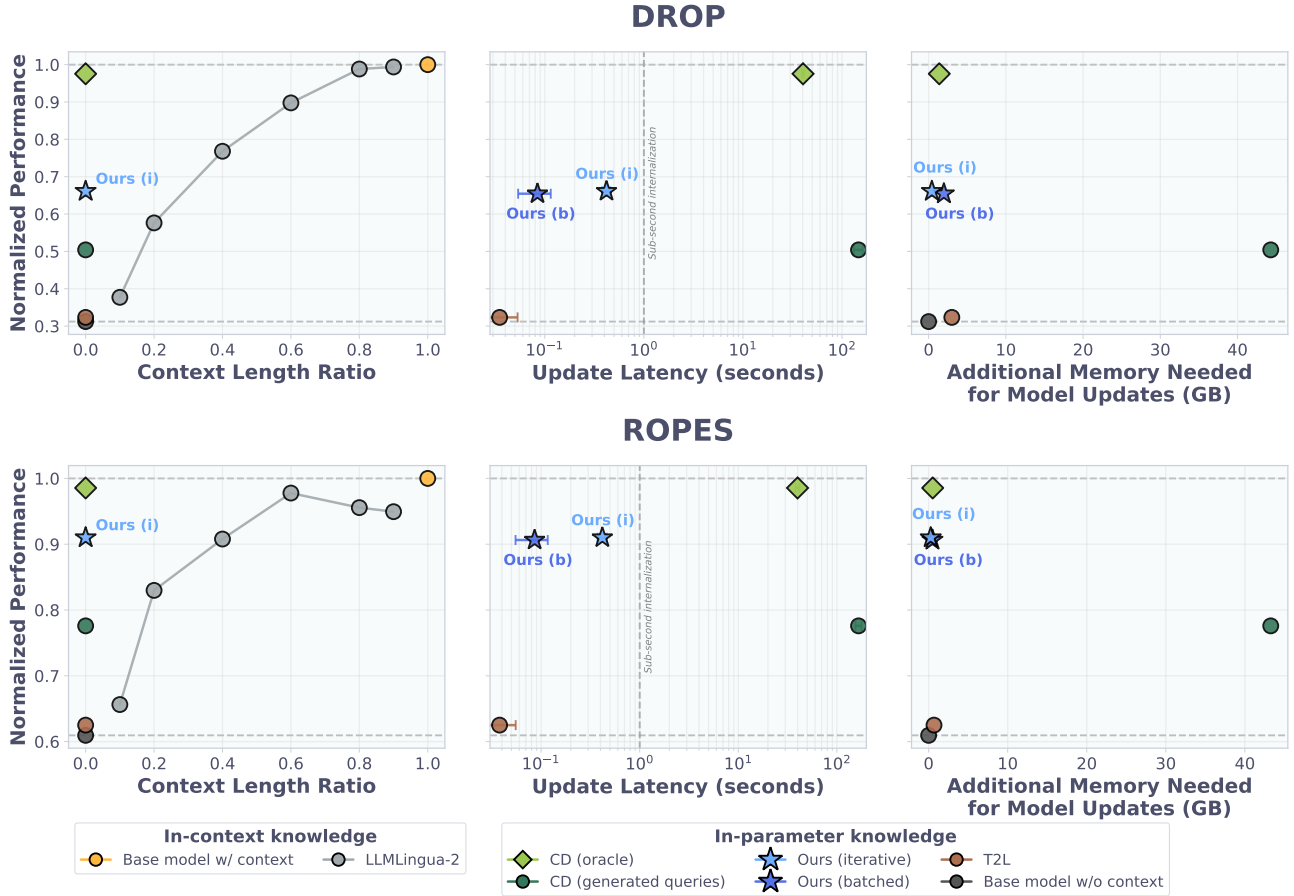
Figure 7. QA performance on DROP and ROPES of all methods compared to used context length ratio (left), update latency (middle), and peak memory used by model updates (right).

Table 10. Performance, update memory, and latency of in-parameter knowledge methods on QASPER benchmark.

Method	Rel. Perf vs Truncated ICL (↑)	Peak Update Memory (GB, ↓)	Mean Update Latency (s, ↓)
CD (oracle query)	0.945	7.665	41.047 ± 1.124
D2L (batched)	0.534	31.153	0.179 ± 0.105
D2L (iterative)	0.543	4.781	0.502 ± 0.097
CD (25 generated queries)	0.480	46.383	410.053 ± 47.922
CD (5 generated queries)	0.446	39.234	87.206 ± 5.053

Table 11. Raw benchmark performance (ROUGE-L F1 score) of the base model with truncated ICL

	SQuAD	DROP	ROPES	2WikiMultihopQA	MultiFieldQA	QASPER
Base model w/ (truncated) context	0.8692	0.4541	0.7457	0.3387	0.3938	0.3839

Passage
During the 1970s and sometimes later, Western and pro-Western governments often supported sometimes fledgling Islamists and Islamist groups that later came to be seen as dangerous enemies. Islamists were considered by Western governments bulwarks against–what were thought to be at the time–more dangerous leftist/communist/nationalist insurgents/opposition, which Islamists were correctly seen as opposing. The US spent billions of dollars to aid the mujahideen Muslim Afghanistan enemies of the Soviet Union, and non-Afghan veterans of the war returned home with their prestige, "experience, ideology, and weapons", and had considerable impact.

Query
Answer the following question. Output only the answer and do not output any other words.
Question: What did the non-Afghan veterans returning home have in addition to their prestige?

Model with internalized passage

Experience, ideology, and weapons.
(ROUGE-L F1 score = 1.0)

Model with internalized query

The provided text states that non-Afghan veterans returned home with their **prestige, "experience," and weapons**.
The text also states that the non-Afghan veterans had a considerable impact.
The text does not provide information about the ideology of the non-Afghan veterans.
(ROUGE-L F1 score = 0.20)

Passage
In 1979, the Soviet Union deployed its 40th Army into Afghanistan, attempting to suppress an Islamic rebellion against an allied Marxist regime in the Afghan Civil War. The conflict, pitting indigenous impoverished Muslims (mujahideen) against an anti-religious superpower, galvanized thousands of Muslims around the world to send aid and sometimes to go themselves to fight for their faith. Leading this pan-Islamic effort was Palestinian sheikh Abdullah Yusuf Azzam. While the military effectiveness of these \"Afghan Arabs\" was marginal, an estimated 16,000 to 35,000 Muslim volunteers came from around the world came to fight in Afghanistan.

Query
Answer the following question. Output only the answer and do not output any other words.
Question: How many Muslims came from around the world to fight in Afghanistan?

Model with internalized passage

16,000 to 35,000
(ROUGE-L F1 score = 1.0)

Model with internalized query

The provided text states that "an estimated **16,000 to 35,000** Muslim volunteers came from around the world came to fight in Afghanistan."
(ROUGE-L F1 score = 0.385)

Passage
The Panthers finished the regular season with a 15\u20131 record, and quarterback Cam Newton was named the NFL Most Valuable Player (MVP). They defeated the Arizona Cardinals 49\u201315 in the NFC Championship Game and advanced to their second Super Bowl appearance since the franchise was founded in 1995. The Broncos finished the regular season with a 12\u20134 record, and denied the New England Patriots a chance to defend their title from Super Bowl XLIX by defeating them 20\u201318 in the AFC Championship Game. They joined the Patriots, Dallas Cowboys, and Pittsburgh Steelers as one of four teams that have made eight appearances in the Super Bowl.

Query
Answer the following question. Output only the answer and do not output any other words.
Question: What team did the Panthers defeat?

Model with internalized passage

Arizona Cardinals
(ROUGE-L F1 score = 1.0)

Model with internalized query

The Panthers defeated the **Cardinals.**
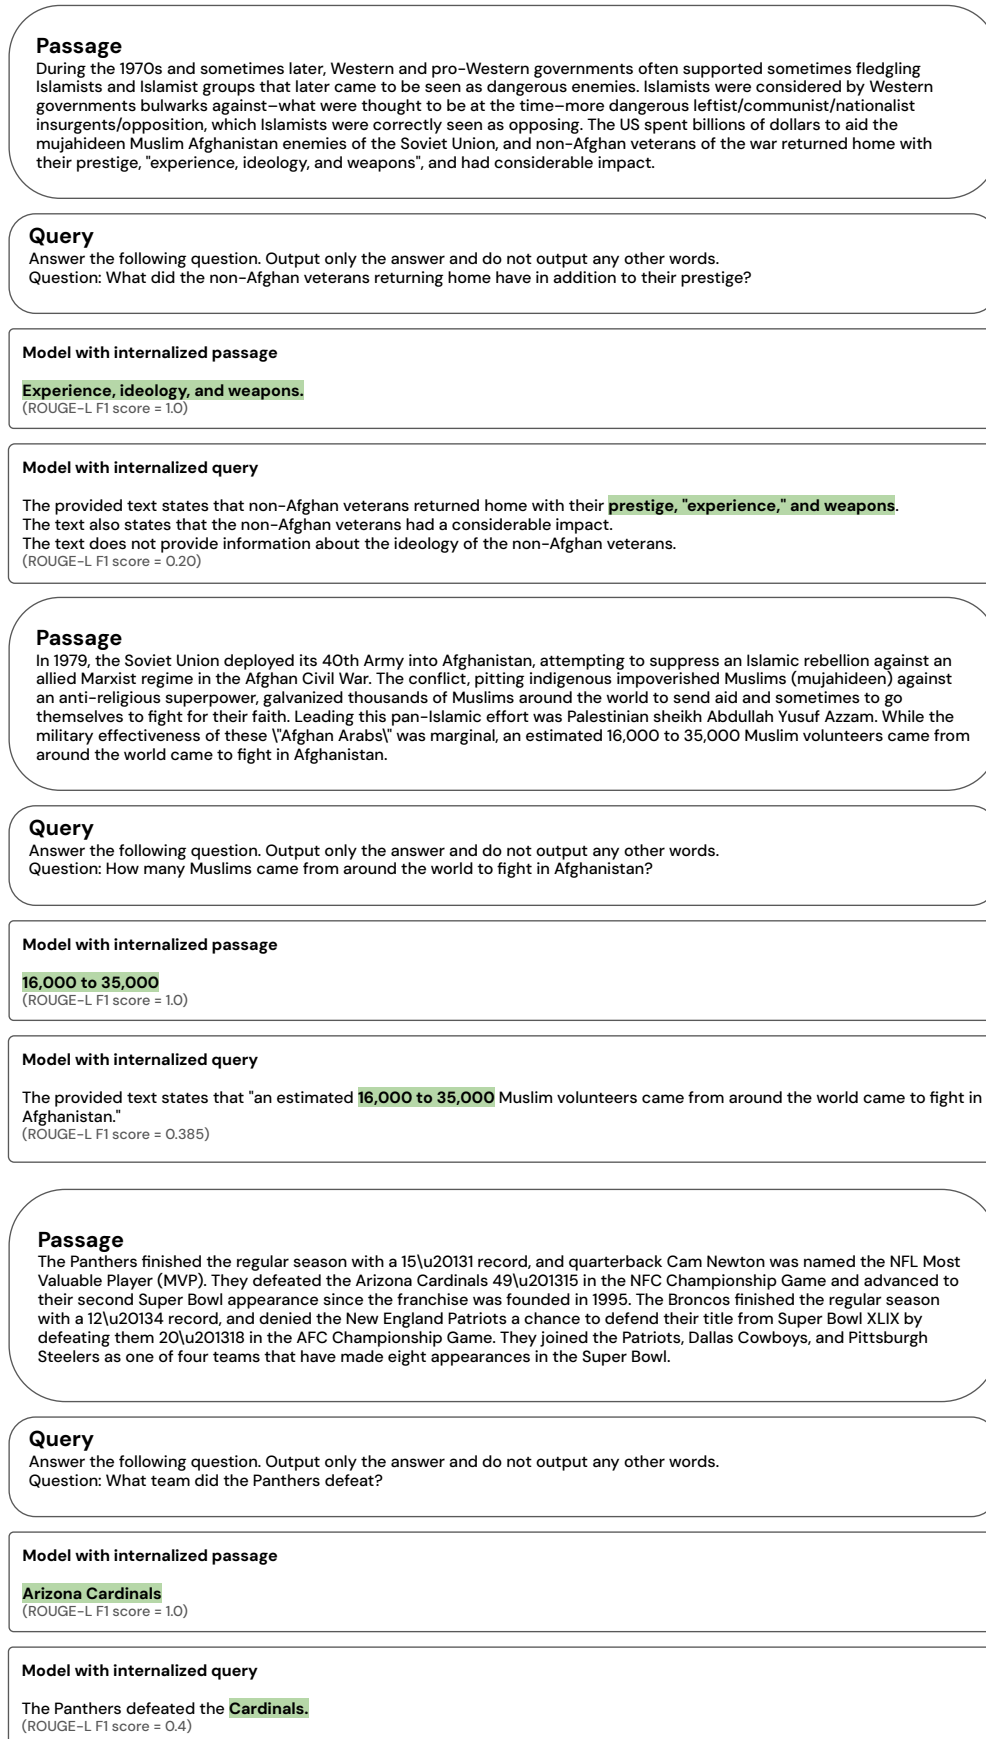(ROUGE-L F1 score = 0.4)

17

Figure 8. Qualitative Results on SQuAD. Green highlight represents words that overlap with the ground-truth labels. The base model with internalized queries sometimes answer correctly but the responses might be verbose, which leads to significantly lower ROUGE-L F1 scores.

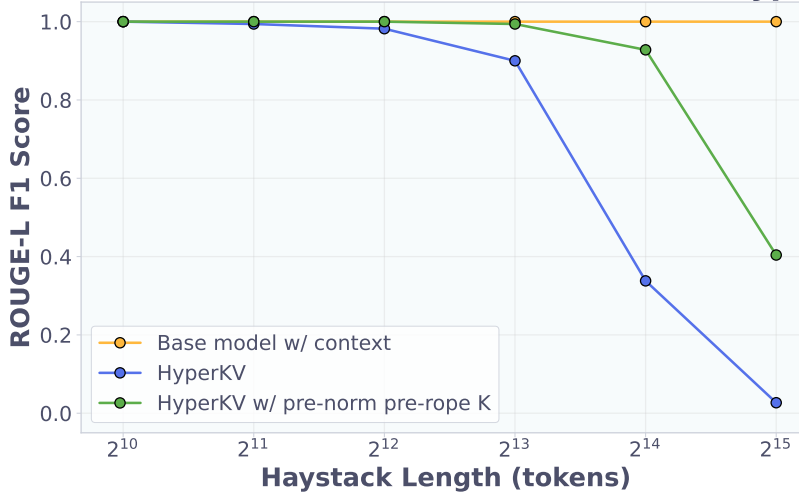**NIAH Retrieval Performance (QWEN-4B + HyperKV)**

Figure 9. NIAH retrieval performance with generated KV cache using Qwen3-4B as the target model.

D. Generating KV cache Instead of LoRA (`Qwen3-4B-Instruct-2507`)

This section aims to explore whether the proposed architecture can be used to generated other kinds of fine-tuning parameterization beyond LoRA. Notably, Eyuboglu et al. (2025) show that prefix-tuning (i.e., KV cache) can be very effective at `CD`. While this parameterization does not internalize the information into the model's weights, it can still compress the KV cache significantly. The hypernetwork in this experiment outputs an equivalent of KV cache of 20 prefix tokens for the target model. We use the same dataset as presented in Section 4. We note that Qwen3-4B has 128K context window, making retrieving a single needle trivial under the haystack lengths we use for the NIAH task. Nonetheless, since the hypernetwork is trained only up to 256 tokens and 8 chunks, exhibiting good performance beyond the training range still indicates strong generalization of the hypernetwork. Similar to Section 4, the input is processed into chunks of 1024 tokens if it is longer than 1024 tokens.

Results: The modified hypernetwork still shows impressive generalization retrieval accuracy beyond its training lengths. Specifically, it achieves near-perfect accuracy up to 8K tokens, beyond which performance degrades gracefully. Interestingly, we observe that the performance degrades quickly if the hypernetwork learns to generate the keys directly: The performance degradation starts around 4K tokens (Figure 9, blue line). Instead, if the generated keys is passed through the key normalization layer and the Rotary Position Embedding (RoPE, Su et al., 2024) layer, the accuracy can be sustained for much longer (Figure 9, green line). We hypothesize that, when the hypernetwork learns to directly output the keys, it has to implicitly emulate RoPE application. Because it has been trained up to 8 chunks, the performance is expected to drop sharply beyond 8 chunks, which is equivalent to 8K tokens in this experiment. The result indeed agrees with this hypothesis as the performance drops sharply beyond 8K tokens when the hypernetwork directly outputs the keys.

Table 12. QA performance on SQuAD with `Mistral-7B-Instruct-v0.2` as the base model. **D2L** maintains the response characteristic of the base model (similar precision and F1 scores) and high factual recall. Generative Adapter, although achieves a higher F1 score, is factually less accurate indicated by the lower recall score.

	ROUGE-L Recall	ROUGE-L Precision	ROUGE-L F1
`Mistral-7B-Instruct-v0.2`	0.919	0.443	0.519
D2L	0.835	0.447	0.515
Generative Adapter	0.643	0.652	0.632

E. Results with Other Models

To test generality of the overall proposed method, we train **D2L** with `Mistral-7B-Instruct-v0.2` and `Qwen3-4B-Instruct-2507` as the base LLMs.

Additionally, the responses from `Mistral-7B-Instruct-v0.2` and `Qwen3-4B-Instruct-2507` are verbose despite instructing the model to output only the answer, e.g., "Answer the following question. Output only the answer and do not output any other words." This behavior makes the ROUGE-L precision and F1 scores much lower compared to `gemma-2-2b-it`. Therefore, for the `Mistral-7B-Instruct-v0.2` and `Qwen3-4B-Instruct-2507` model, we report the ROUGE-L recall score instead of ROUGE-L F1 score. Specific to `Mistral-7B-Instruct-v0.2`, we include Generative Adapter as an additional baseline using the official checkpoint provided by Chen et al. (2025). We also note that Generative Adapter achieves a significantly higher ROUGE-L F1 score than the base model (see Table 12) because of its training method that directly trains the model to output the ground truth tokens via the SFT loss. However, it achieves much lower ROUGE-L recall than that of **D2L** and the base model, indicating that the improved F1 scores comes solely from the fact that Generative Adapter outputs much shorter responses despite answering with less factual accuracy (represented by the ROUGE-L recall score).

The results are presented in Figures 10 to 13. The overall finding is similar to the results collected using `gemma-2-2b-it` as the base model. **D2L** effectively internalizes new information, outperforming all the in-parameter baselines with excellent speed and memory efficiency. The results across model families and sizes indicate the robustness and generality of the proposed method.

We observe that all internalization methods struggle to update long context information for `Mistral-7B-Instruct-v0.2`. Thus, the performance of in-parameter baselines are similar to the no context baseline in this experiment. We offer no explanation as to why this LLM is harder to internalize new long-context knowledge. Investigation of such topic is beyond the scope of this work and we leave the study of such phenomenon for future work.
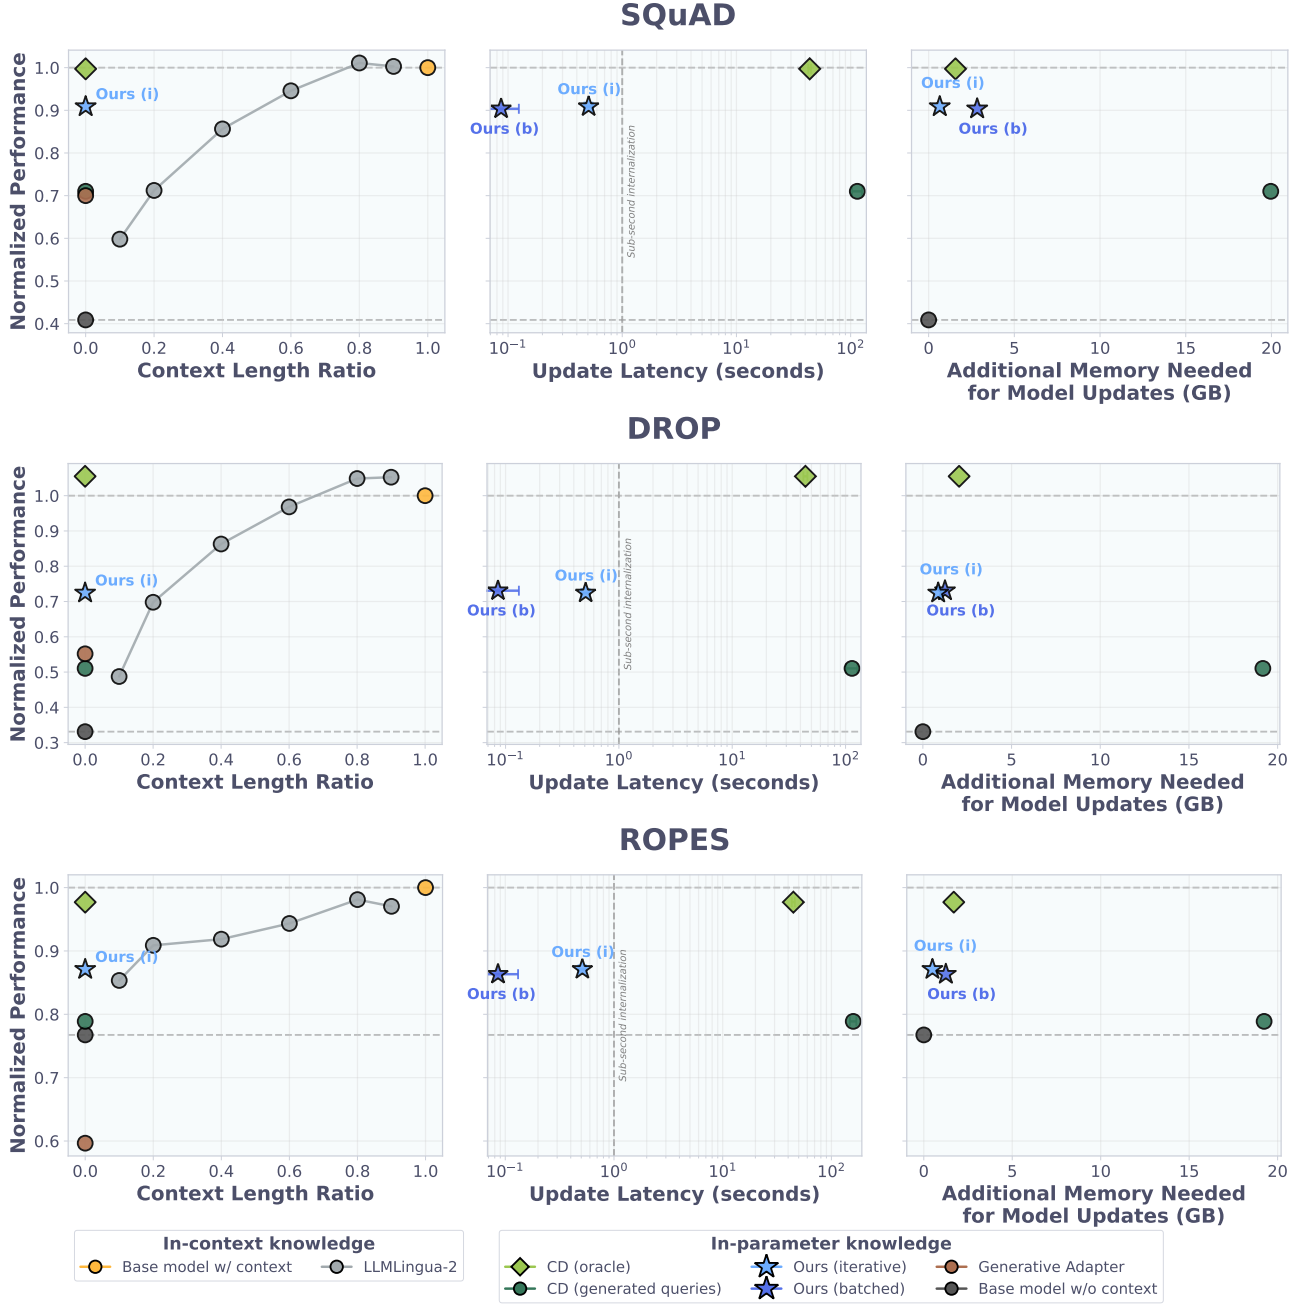
Figure 10. [**D2L** + `Mistral-7B-Instruct-v0.2`] QA performance on DROP and ROPES of all methods compared to used context length ratio (left), update latency (middle), and peak memory used by model updates (right).

Figure 11. [**D2L** + `Mistral-7B-Instruct-v0.2`] Long document QA performance. LLMLingua-2 compresses the input with $[20\%, 40\%, 60\%, 80\%, 90\%]$ compression rates from right to left (gray dots).
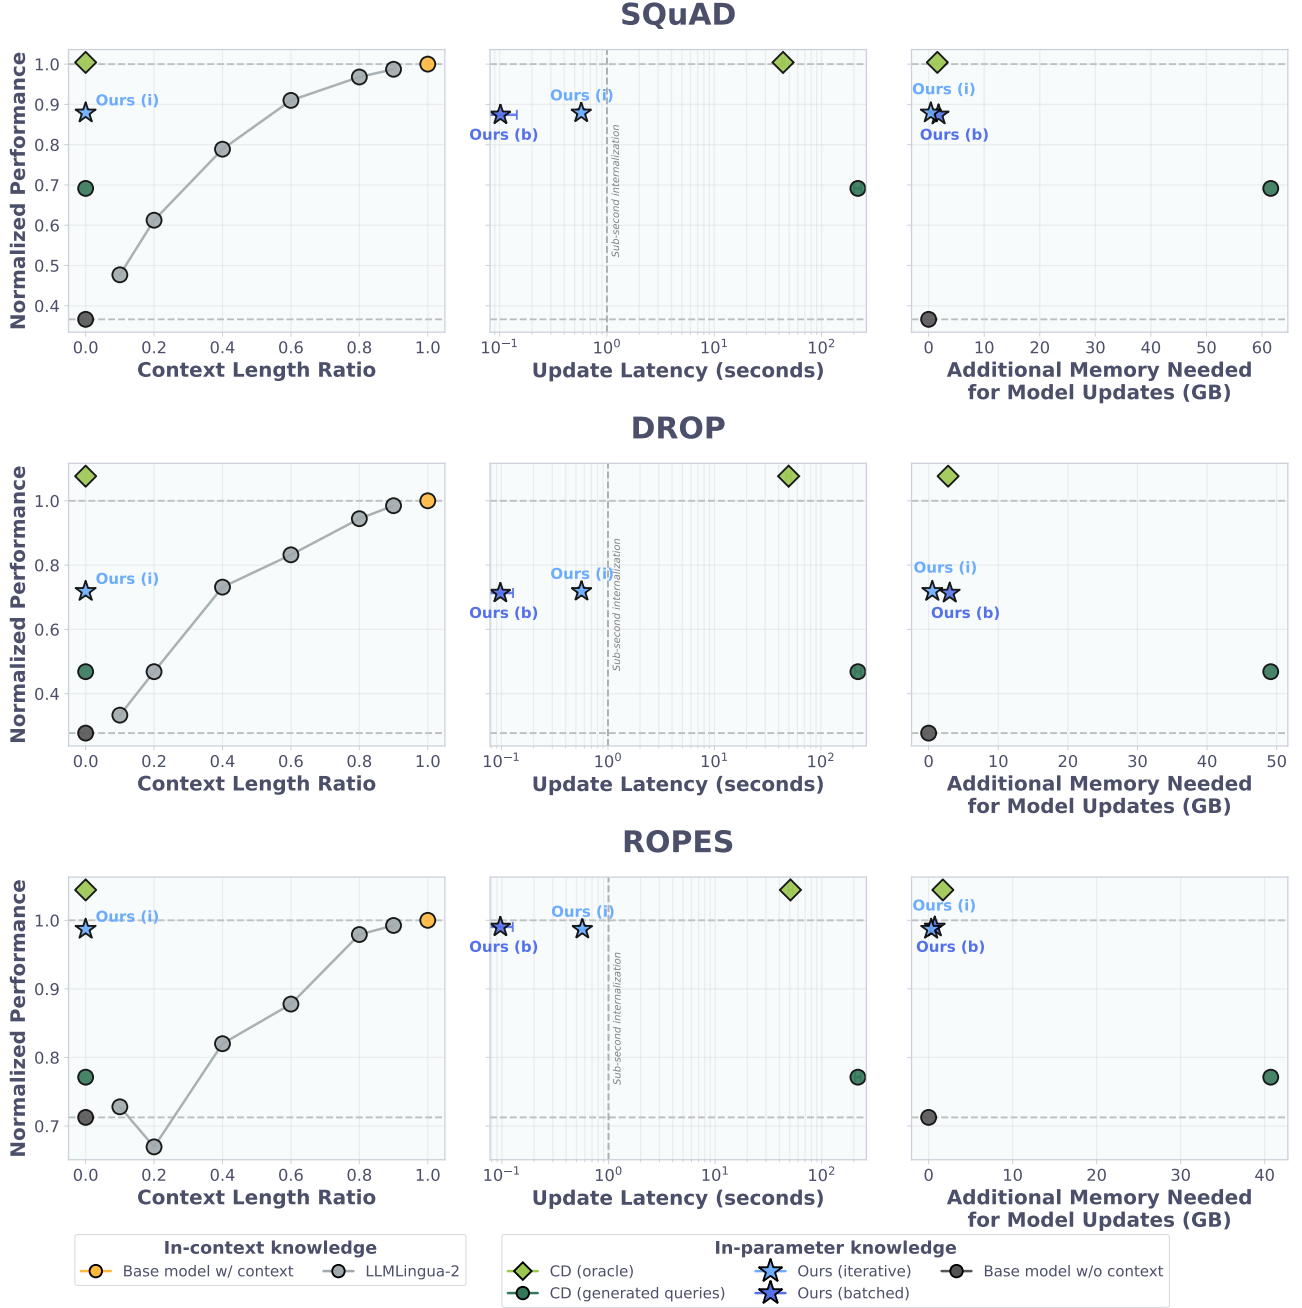
Figure 12. [**D2L** + `Qwen3-4B-Instruct-2507`] QA performance on DROP and ROPES of all methods compared to used context length ratio (left), update latency (middle), and peak memory used by model updates (right).

Figure 13. [**D2L** + `Qwen3-4B-Instruct-2507`] Long document QA performance. LLMLingua-2 compresses the input with $[20\%, 40\%, 60\%, 80\%, 90\%]$ compression rates from right to left (gray dots).

F. Related Work (Extended)

Context Distillation (`CD`): `CD` is a generic and robust method that has been shown to be effective at internalizing value alignment (Askell et al., 2021), new knowledge (Padmanabhan et al., 2023; Qi et al., 2025), long documents (Caccia et al., 2025; Eyuboglu et al., 2025; Kujanpää et al., 2025), conversations (Choi et al., 2023), instructions (Shin et al., 2025), few-shot examples and reasoning traces (Snell et al., 2023; Bhargava et al., 2024). `D2L` aims to learn an approximation of the `CD` process while removing the expensive overheads of `CD`.

Hypernetworks for Meta-Learning: Beyond LLMs, hypernetworks have a long history as meta-learners that amortize task adaptation by directly predicting parameters (von Oswald et al., 2020; Zhao et al., 2020). These methods bypass inner-loop gradient descent at test time by learning a mapping from task embedding to model weights. `D2L` adopts this principle for learning approximate `CD` for LLMs: the hypernetwork maps context information directly to parameter deltas, bypassing the overheads associated with `CD`.

Prompt Compression: A line of work compresses context tokens into condensed soft tokens, reducing the inference cost of LLMs on long sequences (Mu et al., 2024; Pan et al., 2024; Chevalier et al., 2023; Zhang et al., 2025a). These techniques operate directly in the token space (e.g., prompts) by reducing the number of tokens while maintaining performance. `D2L` differs by operating in parameter space via a hypernetwork that predicts weight deltas, which yields persistent, re-usable adaptations.

G. Conclusion

Summary: We propose `D2L`, a hypernetwork that emulates the context distillation (`CD`) process. `D2L` provides instant and inexpensive knowledge internalization through a single forward pass of the hypernetwork by amortizing expensive query generation and backpropagation processes into the meta-training phase (Section 3). On a synthetic NIAH task, `D2L` successfully learns to internalize needle information and effectively extends the context window of the base model to more than $4\times$ its original length (Section 4). On real-world QA tasks, `D2L` outperforms traditional `CD`, with limited query budgets, while significantly reducing memory usage and latency of the internalization process (Section 5.1). Furthermore, we show that `D2L` can zero-shot internalize visual information from a VLM context encoder, allowing the target LLM to have basic visual understanding purely through internalized information (Section 5.2).

Limitations: While leading to significant cost savings compared to `CD`, `D2L` still needs a single expensive meta-training phase. Specifically, an entire training run for meta-learning `CD` for gemma-2-2b-it takes around 5 days on 8 H200 GPUs. Furthermore, in its current form, `D2L` requires retraining the hypernetwork for a new target LLM. We believe that significantly improving the efficiency of the meta-training phase is a fruitful research direction for future work. In general, there is a performance gap between ICL and in-parameter knowledge methods, including `D2L`. Future work could explore more performant methods for internalization, even with update overheads. One such possible approach is combining meta-learned `D2L` with `CD`, using the outputs of `D2L` as the starting point which would be finetuned further by `CD`. This work focuses only on the LoRA parameterization. There are likely better parameterizations that could be more efficient, improve performance, or avoid catastrophic forgetting (Eyuboglu et al., 2025; Lin et al., 2025)

Discussion and Future work: Given that `D2L` can internalize contexts with sub-second latency, we believe that `D2L` has strong implications for inference-time training techniques (Wang et al., 2024; Cao et al., 2025), which still mostly rely on backpropagation-based updates. This capability is also relevant to continual learning (Shi et al., 2025) and personalization (Zhang et al., 2025b), where the base model must iteratively incorporate new knowledge, evolving user preferences, and chat history. Furthermore, `D2L` could help shed light on machine unlearning (Bourtoule et al., 2021) and weight-space interpretability (Olah et al., 2025) by studying and manipulating internalized weights.

H. LLMs Usage Disclosure

In this paper, we use LLMs for early drafting. We also use LLMs for improving writing throughout the paper. Nonetheless, we have verified and will take full responsibility for the contents in this paper.

```
(perceiver): Idefics2Perceiver(
    (modality_projection): Idefics2MLP(
        (gate_proj): Linear(in_features=2304, out_features=9216, bias=False)
        (up_proj): Linear(in_features=2304, out_features=9216, bias=False)
        (down_proj): Linear(in_features=9216, out_features=512, bias=False)
        (act_fn): SiLU()
    )
    (encoder): Idefics2PerceiverResampler(
        (layers): ModuleList(
            (0-7): 8 x Idefics2PerceiverLayer(
                (input_latents_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (input_context_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (cross_attn): Idefics2PerceiverFlashAttention2(
                    (q_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (k_proj): Linear(in_features=512, out_features=512, bias=False)
                    (v_proj): Linear(in_features=512, out_features=512, bias=False)
                    (o_proj): Linear(in_features=2048, out_features=512, bias=False)
                )
                (post_attention_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (pre_ff_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (post_ff_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (mlp): Idefics2MLP(
                    (gate_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (up_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (down_proj): Linear(in_features=2048, out_features=512, bias=False)
                    (act_fn): SiLU()
                )
            )
        )
        (layernorm): Idefics2RMSNorm((512,), eps=1e-06)
    )
    (decoder): Idefics2PerceiverResampler(
        (layers): ModuleList(
            (0): Idefics2PerceiverLayer(
                (input_latents_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (input_context_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (cross_attn): Idefics2PerceiverFlashAttention2(
                    (q_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (k_proj): Linear(in_features=512, out_features=512, bias=False)
                    (v_proj): Linear(in_features=512, out_features=512, bias=False)
                    (o_proj): Linear(in_features=2048, out_features=512, bias=False)
                )
                (post_attention_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (pre_ff_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (post_ff_layernorm): Idefics2RMSNorm((512,), eps=1e-06)
                (mlp): Idefics2MLP(
                    (gate_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (up_proj): Linear(in_features=512, out_features=2048, bias=False)
                    (down_proj): Linear(in_features=2048, out_features=512, bias=False)
                    (act_fn): SiLU()
                )
            )
        )
        (layernorm): Idefics2RMSNorm((512,), eps=1e-06)
    )
)
```

Listing 1. Compact description of the architecture of the Perceiver module.

```
(scaler_A): ParameterDict(  (down_proj): Parameter containing: [torch.cuda.FloatTensor of size
↪ 1x26x8x1 (cuda:5)])
(scaler_B): ParameterDict(  (down_proj): Parameter containing: [torch.cuda.FloatTensor of size
↪ 1x26x8x1 (cuda:5)])
(head): EinMix('bs n_layers n_modules r d_latent -> bs n_layers n_modules r d_lora', 'n_layers
↪ d_latent d_lora', n_layers=26, d_latent=512, r=8, d_lora=11520)
```

Listing 2. Compact description of the hypernetwork's head.

```python
def forward(LLM, hypernet, ctx_ids, ctx_attn_mask, input_ids, input_attn_mask):

    # 1) Encode ctx_ids [n_chunks, seq_len] -> features (Z) [n_chunks, n_layers, seq_len, d]
    features = LLM.forward(ctx_ids, ctx_attn_mask).detach()

    # 2) Hypernet: perceiver
    # [n_chunks, n_layers, r, d_latent]
    emb = hypernet.perceiver(features, ctx_attn_mask)

    # 3) Hypernet: head
    # [n_chunks, n_layers, r, d_in + d_out]
    lora_flat = hypernet.head(emb)

    # 4) Combine across context chunks
    # [n_layers, r * n_chunks, d_in + d_out]
    lora = combine_lora(lora_flat, ctx_ids.shape[0])

    # 5) Apply LoRA to base model layers, then run base model
    apply_lora_to_layers(LLM, lora)
    return LLM.forward(input_ids, input_attn_mask)
```

Listing 3. Psuedocode for the forward pass of the internalized base model.

```python
PROMPT_TEMPLATE = (
    "You are a creative and helpful assistant.\n"
    "You are tasked with generating questions for reading comprehension tests.\n"
    "You will be given a context and you need to generate questions and corresponding answers from
    ↪   the given context.\n"
    "The questions should be highly specific to the information provided in the context, not general
    ↪   questions that suit any context.\n"
    "**DO NOT** hallucinate or make up information.\n\n"
    "### Instructions ###\n"
    "Generate questions and corresponding answers from the given context. The questions should be
    ↪   highly specific to the "
    "information provided in the context, not general questions that suit any context.\n\n"
    "### Context ###\n"
    "{context}\n\n\n"
    "### Rules ###\n"
    "Rules to follow when generating the questions:\n"
    "1. The questions must be specific to the given context and fully answerable from information
    ↪   present in the given context.\n"
    "2. Ask questions that are fact-seeking based on the information provided.\n"
    "3. Make sure the questions are clear and unambiguous.\n"
    "4. Phrases like 'based on the provided context', 'according to the context', 'in the context',
    ↪   etc., are **NOT ALLOWED** to appear in "
    "the questions.\n"
    "5. The questions should not overlap. They should be diverse, covering many aspects of the
    ↪   context.\n"
    "6. Do not give away too much information in the questions. For example, ask 'Who is X?' instead
    ↪   of 'Who is X that did Y?' when Y is clear from the context.\n"
    "7. Ignore the text formatting of the context, e.g., bold, italic, underline, etc.\n"
    "8. Ignore typos, spacing, and grammatical errors in the context.\n\n"
    "Rules to follow when generating the answers:\n"
    "1. The answers must use the (implied) information provided in the context.\n"
    "2. Phrases like 'based on the provided context', 'according to the context', 'in the context',
    ↪   etc., are **NOT ALLOWED** to appear in "
    "the answers.\n"
    "3. Do not just copy words from the context. Answer the question in your own words.\n"
    "4. The answers should be detailed and comprehensive. Please include additional specific details
    ↪   from the context.\n\n"
    "Respond with {n_qa_pairs} question-answer pairs.\n"
    "Always use proper grammar and punctuation.\n"
    "Try to use different question forms and styles.\n"
    "Use simple words and make sure that the answers are clear and comprehensive.\n\n"
    "The question-answer pairs should be in the following format:\n"
    "Question 1: {{question_1}}\n"
    "Answer 1: {{answer_1}}\n"
    "Question 2: {{question_2}}\n"
    "Answer 2: {{answer_2}}\n"
    "..."
)
```

Listing 4. Query generation prompt (first iteration).

```
PROMPT_TEMPLATE_REPEAT = (
    "You are a creative and helpful assistant.\n"
    "You are tasked with generating questions for reading comprehension tests.\n"
    "You will be given a context and you need to generate questions and corresponding answers from
↪    the given context.\n"
    "The questions should be highly specific to the information provided in the context, not general
↪    questions that suit any context.\n"
    "**DO NOT** hallucinate or make up information.\n\n"
    "### Instructions ###\n"
    "Generate questions and corresponding answers from the given context. The questions should be
↪    highly specific to the "
    "information provided in the context, not general questions that suit any context.\n\n"
    "### Context ###\n"
    "{context}\n\n\n"
    "### Example Question-Answer Pairs ###\n"
    "{qa_pairs}\n\n\n"
    "### Rules ###\n"
    "Rules to follow when generating the questions:\n"
    "1. The questions must be specific to the given context and fully answerable from information
↪    present in *or* implied from the given context.\n"
    "2. The questions must *not* be redundant with the example questions-answer pairs provided.\n"
    "3. You should prioritize fact-seeking questions. Consider reversal questions, e.g., asking 'What
↪    causes X to happen?' is valid when 'Y causes X' is presented in the context.\n"
    "4. If all the facts in the context are already covered by the provided examples, you must
↪    generate *more complicated* questions that require reasoning beyond simple information
↪    retrieval.\nThis includes asking about information that can be inferred, requiring
↪    synthesizing information from multiple parts of the text, or understanding relationships
↪    between concepts, events, or individuals mentioned in the context. For example, if the
↪    context says 'The Eiffel Tower was completed in 1889 after 2 years of construction', you can
↪    ask 'When did the construction of the Eiffel Tower begin?'. Here's another example: if the
↪    context says 'Alice is Bob's mother. Bob is Charlie's Dad', you can ask 'Who is Charlie's
↪    grandmother?'.\n"
    "5. Phrases like 'based on the provided context', 'according to the context', 'in the context',
↪    etc., are **NOT ALLOWED** to appear in "
    "the questions.\n"
    "6. The questions should not overlap. They should be diverse, covering many aspects of the
↪    context.\n"
    "7. Do not give away too much information in the questions. For example, ask 'Who is X?' instead
↪    of 'Who is X that did Y?' when Y is clear from the context.\n"
    "8. Ignore the text formatting of the context, e.g., bold, italic, underline, etc.\n"
    "9. Ignore typos, spacing, and grammatical errors in the context.\n\n"
    "Rules to follow when generating the answers:\n"
    "1. The answers must use the (implied) information provided in the context.\n"
    "2. Phrases like 'based on the provided context', 'according to the context', 'in the context',
↪    etc., are **NOT ALLOWED** to appear in "
    "the answers.\n"
    "3. Do not just copy words from the context. Answer the question in your own words.\n"
    "4. The answers should be detailed and comprehensive. Please include additional specific details
↪    from the context.\n\n"
    "Respond with {n_qa_pairs} question-answer pairs.\n"
    "Always use proper grammar and punctuation.\n"
    "Try to use different question forms and styles.\n"
    "Use simple words and make sure that the answers are clear and comprehensive.\n\n"
    "The question-answer pairs should be in the following format:\n"
    "Question 1: {{question_1}}\n"
    "Answer 1: {{answer_1}}\n"
    "Question 2: {{question_2}}\n"
    "Answer 2: {{answer_2}}\n"
    "..."
)
```

Listing 5. Query generation prompt (after first iteration).

```
INTX_TEMPLATES = [
    "Answer the question based on the given passages. Only give me the answer and do not output any
    ↪    other words.\n\nQuestion: {input}",
    "Answer without any explanation.\n\nQuestion: {input}",
    "Based on the provided text, what is the answer to the following question? Provide only the
    ↪    answer.\n\nQuestion: {input}",
    "Extract the answer to the question from the text. Be concise. Do not explain.\n\nQuestion:
    ↪    {input}",
    "What is the answer to this question, based on the context? Respond with the answer
    ↪    only.\n\nQuestion: {input}",
    "Provide a direct answer to the question using the given passages. Do not give any
    ↪    explanation.\n\nQuestion: {input}",
    "Answer the question using only information from the provided text. No extra words.\n\nQuestion:
    ↪    {input}",
    "From the passages, answer the question. Just the answer, please.\n\nQuestion: {input}",
    "Give the answer to the question. Do not include any other text.\n\nQuestion: {input}",
    "The answer to the question is in the text. Find it and state it clearly. No need for
    ↪    explanation.\n\nQuestion: {input}",
    "Concisely answer the question based on the text provided. Don't include any other words. Just
    ↪    the answer.\n\nQuestion: {input}",
    "Read the passages and answer the question with the minimal necessary words.\n\nQuestion:
    ↪    {input}",
    "What is the direct response to the question, according to the text? Avoid
    ↪    explanation.\n\nQuestion: {input}",
    "Please provide only the answer to the question, derived from the text.\n\nQuestion: {input}",
    "Using the provided context, answer the question. Output the answer and nothing else.\n\nQuestion:
    ↪    {input}",
    "Identify the answer in the text and present it without elaboration.\n\nQuestion: {input}",
    "Answer the following question based on the text. Your answer should be brief and to the point.
    ↪    No explanation.\n\nQuestion: {input}",
    "Based on the information given, what is the answer to the question? Only state the
    ↪    answer.\n\nQuestion: {input}",
    "Find the answer to the question in the provided passages and write it down. No
    ↪    explanations.\n\nQuestion: {input}",
    "The question is: {input}. Provide the answer based on the text, and nothing more.",
    "Question: {input}\nAnswer directly based on the text provided. No extra words.",
    "Question: {input}\nPlease provide the answer based on the text. No explanation is needed.",
]
```

Listing 6. Instruction templates for augmenting generated queries.

```
SELF_RESPONSE_TEMPLATE = (
    "You are an honest and helpful assistant.\n\n"
    "# Provided Information\n"
    "{context}\n\n---\n\n"
    "# System Instruction\n"
    "- The information provided is up-to-date information and/or the user instruction.\n"
    "- When the provided information is not relevant to the question, ***ignore*** it and answer the
    ↪    question based on your knowledge.\n"
    "- If the provided information is related to the question, incorporate it in your response.\n"
    "- If the provided information is an instruction, follow the instruction carefully.\n"
    "\n---\n\n"
    "# User Input\n"
    "{question}"
)
```

Listing 7. Prompt template for generating self-responses.

```
QA_PROMPT = (
    "Answer the following question. Output only the answer and do not output any other
    ↪    words.\n\nQuestion: {input}"
)

QASPER_QA_PROMPT = (
    'Answer the question as concisely as you can, using a single phrase or sentence if possible.\nIf
    ↪    the question cannot be answered based on the information in the article, write
    ↪    "unanswerable".\nIf the question is a yes/no question, answer "yes", "no", or "unanswerable".
    ↪    Do not provide any explanation.\n\nQuestion: {input}'
)
```

Listing 8. Evaluation prompts.